

Performance Analysis

CS6133 - Computer Architecture I

Vikram Padman

Polytechnic Institute of New York University

vikram@poly.edu

- ➊ Introduction
 - Limitations
- ➋ Performance Analysis
- ➌ Speedup
 - Amdahl's Law
- ➍ Dependability
- ➎ Activities

- Performance is a quantitative measure of how well a system works
- Why is performance important?
 - To compare and contrast computing system. Which system has the best performance for the price
 - For trade studies: Evaluate a system for a given application (its usefulness)
 - To understand a system's underlying architecture
 - To forecast future needs and estimate improvement cost

- A Car's Performance metrics:
 - Engine's horse power
 - 0-60Mph and/or 1/4 mile time
 - Miles per gallon
 - Braking distance
 - Seating and cargo capacity
- A computer's performance metrics:
 - Power
 - Clock speed
 - Battery life
 - Execution time / Responsiveness
 - Storage capacity

Performance since 1980's

Performance
Analysis

Vikram
Padman

Agenda

Introduction

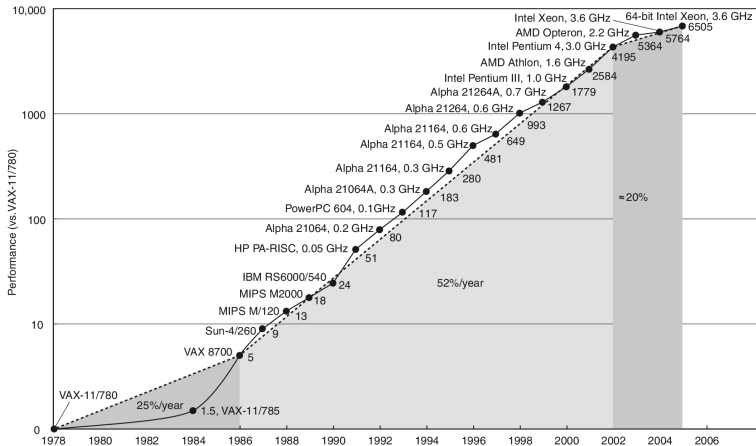
Limitations

Performance

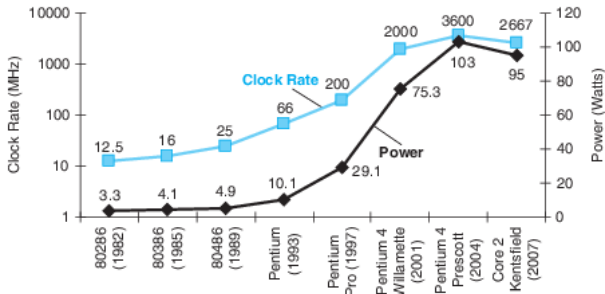
Speedup

Dependability

Activity



The Power Wall



- Power dissipation is an issue in today's technology.
- Smaller transistors could switch faster and more could be packed in a given area
- Smaller transistors \Rightarrow smaller wires. Smaller wires increase resistance and capacitance.

- **Dynamic Power:** Power consumed by a active transistor
- **Static Power:** Quiescent or leakage power consumed by an idle transistors
- $Power_{dynamic} = 1/2 * C * V^2 * T_r$
- $Power_{static} = I_{leakage} * V$
- C = Total (gate + wire) capacitance per transistor
- V = Rail voltage, T_r = Toggle rate, I = Current
- Total power for a chip:

$$Power_{total} = (Power_{dynamic} + Power_{static}) * N_t$$
- N_t = Number of transistors

The Scaling Wall

Performance
Analysis

Vikram
Padman

Agenda

Introduction

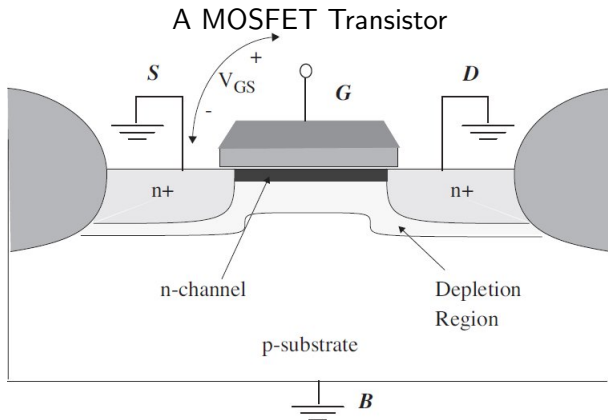
Limitations

Performance

Speedup

Dependability

Activity



- Reduction in transistor's size increases leakage current
- Increase in transistor count increases static power
- Reduction in rail voltage could introduce instability due to process variations
- Increase in frequency increases dynamic power

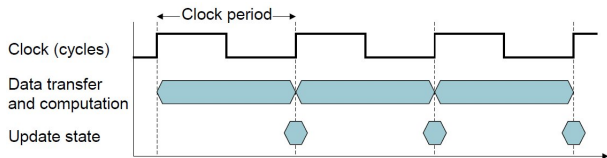
- **Execution Time** = Total run time of a program
 - Time could simply be “wall-clock time” or “elapsed time”
- **Performance** = $1 / (\text{Execution time})$
 - Clock speed, parallelism, I/O speed, and work load determines performance
- **Relative Performance** = Performance of one system compared to another
 - $R_p = \text{Performance}_A / \text{Performance}_B = \text{Execution}_B / \text{Execution}_A$
 - System A is R_p time faster than system B
- R_p accuracy depends on how execution time is measured in each system and its consistency

- Cost-Performance ratio:

$$C_r = \text{Cost} * \text{Execution time or Cost} / \text{Performance}$$

- Let's consider system A and B running application App_1 and App_2
 - Execution time of App_1 in A = 8 and cost of A = 10,
 $C_r = 80$
 - Execution time of App_2 in A = 1 and cost of A = 10,
 $C_r = 10$
 - Execution time of App_1 in B = 2 and cost of A = 15,
 $C_r = 30$
 - Execution time of App_2 in B = 2 and cost of A = 15,
 $C_r = 30$

- CPU's performance is measured by the number of clock cycles it consumes to execute a given program
- Clock rate = cycles per second, measured in Hz , Khz or Ghz
- Clock period = time between rising edge of a clock signal



CPU time consumed by a given program =

- $CPU_{time} = I_{tot} * Cycles_i * Clockrate$
- I_{tot} : Total number of Instructions
 - Depends on CPU's architecture, ISA, compiler, coding style and application
- $Cycles_i$: Clock cycles per instruction
 - Depends on CPU's architecture and ISA
- $Clockrate$:
 - ASIC design, logic design, node size and silicon manufacturing technology

Clock Cycles Per Instruction (CPI)

- **CPI:** Average number of clock cycles per instruction:

$$CPI = C_{el} / I_{tot}$$
 - I_{tot} = Number of instruction in a program. A program could have many types of instruction(s)
 - C_{el} = Number of elapsed clock cycle(s)
- Let consider a program AP_1
 - N = types of instruction in a program AP_1
 - I_t = number of instruction of type t in AP_1

$$\Rightarrow I_{tot} = \sum_{t=1}^N I_t$$
 - CPI_t = clock cycles(s) to execute 1 instruction of type t

$$C_{el} = \sum_{t=1}^N (I_t * CPI_t)$$

$$\Rightarrow CPI = C_{el} / I_{tot} = \sum_{t=1}^N (I_t / I_{tot} * CPI_t)$$

Clock Cycles Per Instruction (CPI)

Continuing with AP_1

- Suppose AP_1 has 100 instructions (I_{tot}) and below is the distribution:

t	Instruction	I_t	CPI_t	$I_t/I_{tot} * CPI_t$
1	load	35	4	$35/100 * 4 = 1.4$
2	store	10	2	$10/100 * 2 = 0.2$
3	mul	15	35	$15/100 * 35 = 5.25$
4	jump	15	10	$15/100 * 10 = 1.5$
5	add	15	1	$20/100 * 1 = 0.2$
6	sub	10	1	$10/100 * 1 = 0.1$
$CPI = \sum_{t=1}^6 (I_t/I_{tot}) * CPI_t$				8.65

- $C_{el} = CPU_{time} * ClockRate = I_{tot} * CPI$
- For example: AP_2 with $10 * 10^7$ instructions takes a second to complete on a $500Mhz$ CPU than
 $CPI = 1 * 500 * 10^6 / (10 * 10^7) = 5$

- **MIPS** A theoretical peak instruction execution rate
- $MIPS = I_{tot} / (CPU_{time} * 10^6)$
 $\Rightarrow MIPS = I_{tot} / (C_{el} * CycleTime * 10^6)$
 $\Rightarrow MIPS = (I_{tot} * ClockRate) / (I_{tot} * CPI * 10^6)$
 $\Rightarrow MIPS = ClockRate / (CPI * 10^6)$
- Using MIPS rating is tied to a particular program running on specific machine. MIPS rating is not a reliable measure that could be used to compare two different systems.
- **DMIPS** Dhrystone MIPS, on the other hand, is a useful measure
- Dhrystone is a benchmark program developed by Reinhold P. Weicker to measure a CPU's integer processing performance. Dhrystone could be obtained here <http://www.netlib.org/benchmark/dhry-c>

- **MFLOPS** A theoretical maximum floating point instruction execution rate.
- MFLOPS are measured using benchmark programs such as:
 - Whetstone : A floating point benchmark software
 - LINPACK : A software library for linear algebra written in fortran
 - LAPACK : Latest version of LINPACK

- **Speedup** is defined by:
$$OriginalCPU_{time} / NewCPU_{time}$$
- Overall improvement of a system's performance could be achieved by improving a subset of the system.
- Most of the performance increase come from improving the common case(s)

- **Amdahl's Law** is used to calculate the overall performance improvement achieved by improving a subset of a system.
- Amdahl's Law states that performance improvement to be gained by using a faster mode is limited by the fraction of time the faster mode can be used.
- $\text{Speedup} = \text{Total performance with faster mode} / \text{Total performance without faster mode}$
- $\text{Speedup} = \text{Execution time without faster mode} / \text{Execution time with faster mode}$

- $Speedup_{overall} = CPU_{Oldtime} / CPU_{Newtime}$
- $\begin{aligned}
 CPU_{Newtime} &= \\
 &UnenhancedCPU_{time} + EnhancedCPU_{time} \\
 &\Rightarrow CPU_{Oldtime} * (1 - F_{enh}) \\
 &\quad + (CPU_{Oldtime} * F_{enh}) / Speedup_{enh} \\
 &\Rightarrow CPU_{Oldtime} * ((1 - F_{enh}) + F_{enh} / Speedup_{enh})
 \end{aligned}$
- $\Rightarrow Speedup_{overall} = 1 / ((1 - F_{enh}) + F_{enh} / Speedup_{enh})$

Implications of Amdahl's Law

Performance
Analysis

Vikram
Padman

Agenda

Introduction

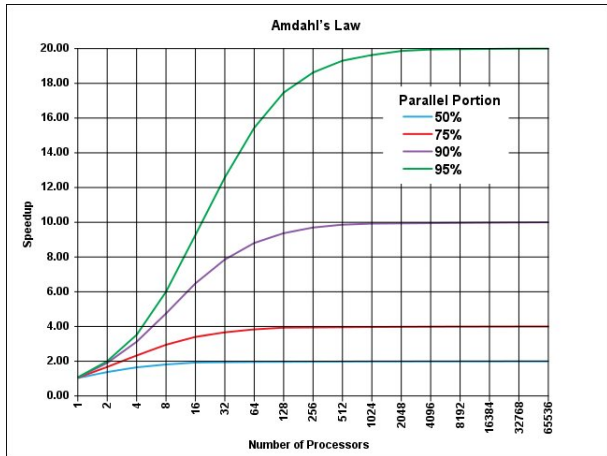
Performance

Speedup

Amdahl's Law

Dependability

Activity



- **Dependability** Today, computers are integral part of our society. It is important to understand the reliability and availability of computers.
- Metrics for dependability:
 - ① **MTTF**: Mean time to failure is a reliability measure
 - ② **MTTR**: Mean time to repair is a service interruption measure
 - ③ **MTBF**: Mean time between failure is a commonly used term which is $= MTTR + MTTF$
 - ④ Module availability $= MTTF / (MTTF + MTTR)$

Read Section 1.7-Dependability in the course text book and answer the following questions:

- ① What is the difference between SLA, SLO, MTTF, MTBF and MTTR? Describe with examples how these metrics are used in real life. (3 pages max)
- ② You are being asked to build a 15TB storage vault with SAS drives + spares to replace ageing SCSI storage vault in your organization. The storage vault you build/buy should last for at least 8 years, performance is crucial, downtime is very expensive and yearly maintenance cost should be minimum.
 - ① Which metric(s) would you use to discriminate drives and drive models made by different manufacturers? and how?
 - ② Pick a SAS drive that would be suitable for the storage vault and support your decision with analysis.
 - ③ Calculate the optimal number of spare drives?
 - ④ Which metric(s) would you use to estimate yearly maintenance cost?

Describe and explain the purpose of the following benchmark standards and software:

- 1 BAPCo
- 2 SPEC
- 3 TPC
- 4 Coremark
- 5 Dhrystone
- 6 Whetstone

Consider the following profiles of two application that are running on *CPU – ZX*:

APP₁ Profile

t	Instruction	I_t	CPI_t
1	load	35	4
2	store	15	2
3	mul	2	35
4	jump	25	10
5	add	35	5
6	sub	15	5

APP₂ Profile

t	Instruction	I_t	CPI_t
1	load	10	4
2	store	20	2
3	mul	4	35
4	jump	25	10
5	add	25	5
6	sub	10	5

Answer the following questions:

- 1 Calculate the CPI, performance and relative performance of *APP₁* & *APP₂*?
- 2 Select three instruction in *CPU – ZX*'s ISA that when improved would provide the most performance increase. Support your answer with thorough analysis.
- 3 Assume that you were able to make CPI_t of the three instruction that you picked to 3. Calculate overall speedup of *APP₁* and *APP₂*.