

# Input/Output Subsystems

## CS6133 - Computer Architecture I

Vikram Padman

Polytechnic Institute of New York University

[vikram@poly.edu](mailto:vikram@poly.edu)

- ➊ Introduction
- ➋ Addressing
- ➌ Methodology
  - Programmed
  - Interrupt Driven
  - Direct Memory Access (DMA)
- ➍ Data Transfer
  - Synchronous
  - Asynchronous
- ➎ Interconnects
  - System or Shared BUS
  - Point-to-Point
    - Parallel
    - Serial
- ➏ Performance

- CPU's interface to the external world
- Plays an important role in overall system's performance
  - Most application interact with external world
  - Has a profound effect on the CPI
- Essentially, there are three types of Input/Output devices
  - **Human Machine Interface - HMI** devices
    - Keyboard, mouse, joystick, monitors, Google Glass ..etc
  - **Secondary Storage** devices
    - SATA/SAS hard drives, flash disk, SDHC card ..etc
  - **Communication** devices
    - Network Interface Cards, WiFi cards, Modems ...etc



## ① I/O Address Space

- A unique address space that could only be accessed by special instruction
- CPU identifies each I/O using a unique address
- Uses “IN” or “OUT” instruction to move data
- Special I/O instruction could only be executed by root user or OS

## ② Memory Mapped I/O

- Uses general memory address that could be accessed by any instructions.
- CPU treats an I/O device as memory and contents of memory mapped I/O device are cache-able, unless stated otherwise
- Access protection could be changed by root or system user

# Input/Output Addressing

## Input/Output Subsystems

Vikram  
Padman

## Agenda

## Introduction

## Addressing

## Methodology

## Data Transfer

## Interconnects

## Performance

## Activity

```

KKV-SU ~ # lspci
00:00.0 Host bridge: Advanced Micro Devices [AMD] nee ATI RD890 PCI to PCI bridge (external gfx0 port A) (rev 02)
00:02.0 PCI bridge: Advanced Micro Devices [AMD] nee ATI RD890 PCI to PCI bridge (PCI express gpp port B)
00:03.0 PCI bridge: Advanced Micro Devices [AMD] nee ATI RD890 PCI to PCI bridge (PCI express gpp port C)
00:09.0 PCI bridge: Advanced Micro Devices [AMD] nee ATI RD890 PCI to PCI bridge (PCI express gpp port H)
00:0a.0 PCI bridge: Advanced Micro Devices [AMD] nee ATI RD890 PCI to PCI bridge (external gfx1 port A)
00:11.0 SATA controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 SATA Controller [AHCI mode]
00:12.0 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 USB OHCI0 Controller
00:12.1 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0 USB OHCI1 Controller
00:12.2 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 USB EHCI Controller
00:13.0 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 USB OHCI0 Controller
00:13.1 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0 USB OHCI1 Controller
00:13.2 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 USB EHCI Controller
00:14.0 SMBus: Advanced Micro Devices [AMD] nee ATI SBx00 SMBus Controller (rev 3d)
00:14.3 ISA bridge: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 LPC host controller
00:14.4 PCI bridge: Advanced Micro Devices [AMD] nee ATI SBx00 PCI to PCI Bridge
00:14.5 USB controller: Advanced Micro Devices [AMD] nee ATI SB7x0/SB8x0/SB9x0 USB OHCI2 Controller
00:18.0 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 0
00:18.1 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 1
00:18.2 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 2
00:18.3 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 3
00:18.4 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 4
00:18.5 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 5
00:19.0 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 0
00:19.1 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 1
00:19.2 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 2
00:19.3 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 3
00:19.4 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 4
00:19.5 Host bridge: Advanced Micro Devices [AMD] Family 15h Processor Function 5
01:00.0 VGA compatible controller: NVIDIA Corporation G96 [GeForce 9500 GT] (rev a1)
02:00.0 RAID bus controller: LSI Logic / Symbios Logic MegaRAID SAS 1078 (rev 04)
03:00.0 Ethernet controller: Intel Corporation 82574L Gigabit Network Connection
04:00.0 Ethernet controller: Intel Corporation 82574L Gigabit Network Connection
05:03.0 Multimedia audio controller: C-Media Electronics Inc CM18788 [Oxygen HD Audio]
    
```

# Input/Output Addressing

## Input/Output Subsystems

Vikram  
Padman

## Agenda

## Introduction

## Addressing

## Methodology

## Data Transfer

## Interconnects

## Performance

## Activity

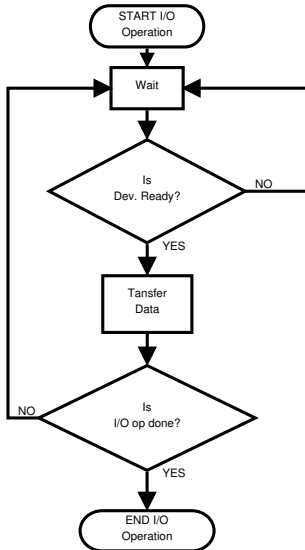
```
01:00.0 VGA compatible controller: NVIDIA Corporation G96 [GeForce 9500 GT] (rev a1) (prog-if 00 [VGA controller])
Subsystem: Device 196e:0643
Flags: bus master, fast devsel, latency 0, IRQ 24
Memory at fb000000 (32-bit, non-prefetchable) [size=16M]
Memory at c0000000 (64-bit, prefetchable) [size=512M]
Memory at fc000000 (64-bit, non-prefetchable) [size=32M]
I/O ports at a800 [size=128]
[virtual] Expansion ROM at faf80000 [disabled] [size=512K]
Capabilities: [60] Power Management version 3
Capabilities: [68] MSI: Enable- Count=1/1 Maskable- 64bit+
Capabilities: [78] Express Endpoint, MSI 00
Capabilities: [b4] Vendor Specific Information: Len=14 <?>
Capabilities: [100] Virtual Channel
Capabilities: [128] Power Budgeting <?>
Capabilities: [600] Vendor Specific Information: ID=0001 Rev=1 Len=024 <?>
Kernel driver in use: nvidia

02:00.0 RAID bus controller: LSI Logic / Symbios Logic MegaRAID SAS 1078 (rev 04)
Subsystem: LSI Logic / Symbios Logic MegaRAID SAS 8888ELP
Flags: bus master, fast devsel, latency 0, IRQ 28
Memory at fe9c0000 (64-bit, non-prefetchable) [size=256K]
I/O ports at b000 [size=256]
Memory at fe980000 (64-bit, non-prefetchable) [size=256K]
Expansion ROM at fe960000 [disabled] [size=128K]
Capabilities: [b0] Express Endpoint, MSI 00
Capabilities: [c4] MSI: Enable- Count=1/4 Maskable- 64bit+
Capabilities: [d4] MSI-X: Enable+ Count=4 Masked-
Capabilities: [e0] Power Management version 2
Capabilities: [ec] Vital Product Data
Capabilities: [100] Power Budgeting <?>
Kernel driver in use: megaraid_sas
```

- Commonly known as Polling I/O
  - ① The CPU queries for the readiness of the I/O
  - ② When the device is ready CPU transfers one or more data units
  - ③ Waits for an acknowledgement
  - ④ Goes back to step 1 to transfer more data
- CPU is in total control of the I/O device
- Busy wait loops are used in code until device is ready
- Many real time systems prefer or require programmed I/O method



# Programmed I/O



- Programmed I/O, in general, is not efficient
- CPU spends considerable amount of time waiting
- CPU is in greater control, so PIO method is preferred in some real time systems

- **Periodic Polling** CPU periodically checks to see if the device is ready for I/O
- **Busy or Spin Wait** CPU stays in a busy wait loop and continuously checks for I/O device's readiness

- Preferred I/O method in modern computing systems
  - ① CPU sends a message to I/O device and proceed to execute a different thread or process
  - ② I/O device interrupts' the CPU when it is ready
  - ③ CPU stops current execution, switches back to I/O process and performs I/O operation
- Traditionally, interrupts are hard wire connecting the CPU and I/O device. Modern systems use message signalled interrupt **MSI**
- An interrupt causes the CPU to jump to interrupt service routine (ISR) to service the I/O device
- Data transfers are done by CPU or using DMA

- Used in conjunction with Interrupt driven I/O method
- A section of memory is allocated to each I/O device
  - ① CPU sends a message to I/O device for data transfer
  - ② I/O device transfers data to/from the memory allocated to it
  - ③ When I/O device completes requested I/O operation, it interrupts CPU to signal completion
- Memory marked for DMA are not cache-able, unless its contents cannot be modified by the device

- There are two methods to transfer data between the CPU and I/O device

## ① Synchronous Data Transfer

- Data is transferred relative to an I/O clock
- I/O clock is a multiple of CPU's internal clock
- Transfer speed and timing is fixed and does not change

## ② Asynchronous Data Transfer

- Dedicated signal(s) are used to control data transfer
- Validity of data is indicated by a "data valid" signal
- Transfer speed and timing could vary

- I/O clock is phase aligned to CPU's clock
  - Clock could be  $1/n * CPU_{clk}$  where  $n$  is a whole number
  - Dedicated wires are used to carry clock signal
  - Transfer protocol is defined with reference to I/O clock signal
- Advantages
  - Very fast, high performance I/O devices are synchronous devices
  - Transfer protocols are simple
- Disadvantages
  - Distance between CPU and I/O must be short due to clock skew
  - In a shared BUS environment all devices must run at the same speed

# Synchronous Data Transfer

## Input/Output Subsystems

Vikram Padman

## Agenda

### Introduction

### Addressing

### Methodology


### Data Transfer

### Synchronous Asynchronous

### Interconnects

### Performance

### Activity



- Clocks are not used
- Transfers are controlled by handshaking protocols using dedicated wires
- Advantages
  - Distances between CPU and I/O device could be large (meters)
  - Devices sharing the physical medium could run at different speed
- Disadvantages
  - Handshaking protocols could be complex
  - Usually slow due to complexity
  - Latencies could be undeterministic



# Asynchronous Data Transfer

## Input/Output Subsystems

Vikram Padman

## Agenda

Introduction

Addressing

Methodology

Data Transfer


Synchronous

Asynchronous

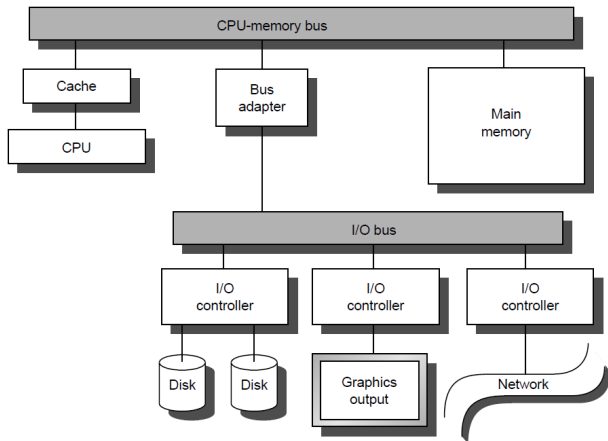
Interconnects

Performance

Activity



- **BUS** A shared communication link between multiple devices
  - Physical wires connecting multiple devices
- Advantages
  - BUS implementations are cheaper
  - Easy to add or remove devices
- Disadvantages
  - Has to be relatively short due to electrical problems
  - Bus shew, crosstalk and dispersion are common
  - BUS contention could arise due to sharing



## Why talk about High Speed Serial Communication?

- It is the future
- Used in all modern computing systems. Example: SAS/SATA, PCIe, Intel's QPI, USB, ThunderBolt, AMD's HyperTransport, SRIO, HDMI ... etc
- Easy to manufacture, scalable and very flexible
- In most cases it uses only four wires to transfer giga bits of data in a second
- Supported by many chip and PCB manufactures.
- Most high speed serial protocols support automatic speed throttling, error correction and on-board diagnostics.

# Parallel vs Serial Communication

## Parallel Communication

Input/Output  
Subsystems

Vikram  
Padman

Agenda

Introduction

Addressing

Methodology

Data Transfer

Interconnects

BUS

P2P


PC vs SC

SERDES

Performance

Activity

- Number of connections
- Data synchronization and alignment
- Resistance to interference
- Chip package size
- Speed limitations



# Parallel vs Serial Communication

## Serial Communication

Input/Output  
Subsystems

Vikram  
Padman

Agenda

Introduction

Addressing

Methodology

Data Transfer

Interconnects

BUS


P2P

PC vs SC

SERDES

Performance

Activity



# SERializer DE-Serializer

## Input/Output Subsystems

Vikram Padman

## Agenda

## Introduction

## Addressing

## Methodology

## Data Transfer

## Interconnects

## BUS

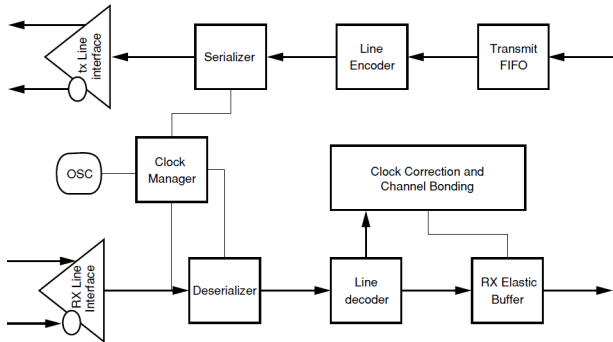
## P2P

## PC vs SC

## SERDES

## Performance

## Activity



- **Throughput / Bandwidth**

- Total data exchanged between the processor/memory and the I/O device within a unit time

- **Latency / Response Time**

- Total elapsed time to transfer one data unit
- An important measure for real time systems where consistent latency is important

- High throughput, low latency I/O device are desired by many applications



Read Chapter 1 to 4 from “High-Speed Serial I/O Made Simple - Xilinx”<sup>1</sup> (on-line book) and answer the following questions

- ① What is LVDS and when should it be used. What are the disadvantages of using LVDS?
- ② What is the caveat in using Clock and Data Recovery(CDR)? Are there any serial protocols that does not use CDR?
- ③ Why is 8b/10b encoding used for serial communication? and what is the overhead in using 8b/10b encoding.
- ④ What is the difference between 8b/10b and 64b/66b encoding?
- ⑤ What is an “Eye Patterns”?
- ⑥ What is a “Scrambler” and which popular serial protocol uses it and why?

<sup>1</sup><http://www.xilinx.com/publications/archives/books/serialio.pdf>