Chapter 2
A Single Segment Network – Data Link Layer

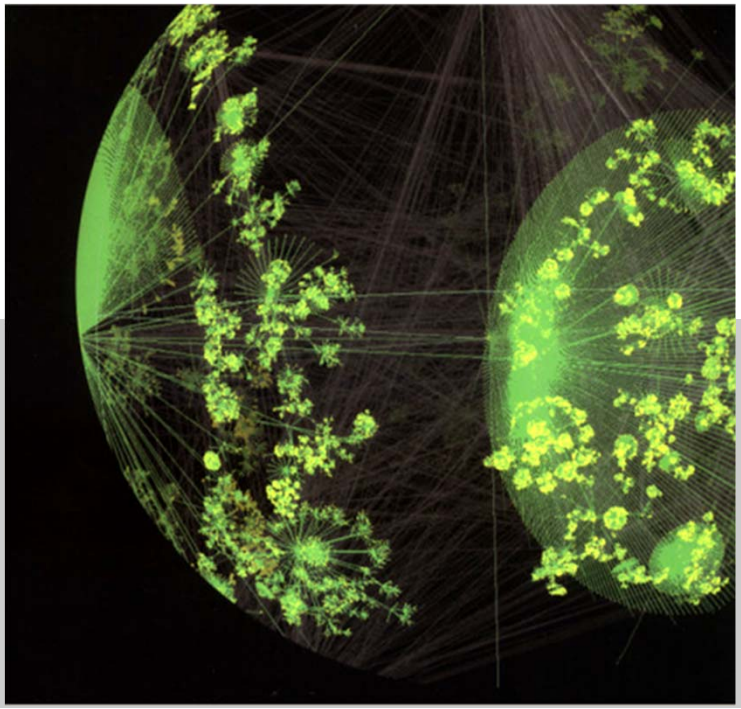

TCP/IP Essentials

A Lab-Based Approach

Spring 2017

Data Link Layer

Main tasks of the data link layer:

- Transfer network layer data from one machine to another machine via “a data link”.
- Convert the data between raw bit stream of the physical layer and groups of bits → bytes → frames.
- Perform flow control between sender and receiver.

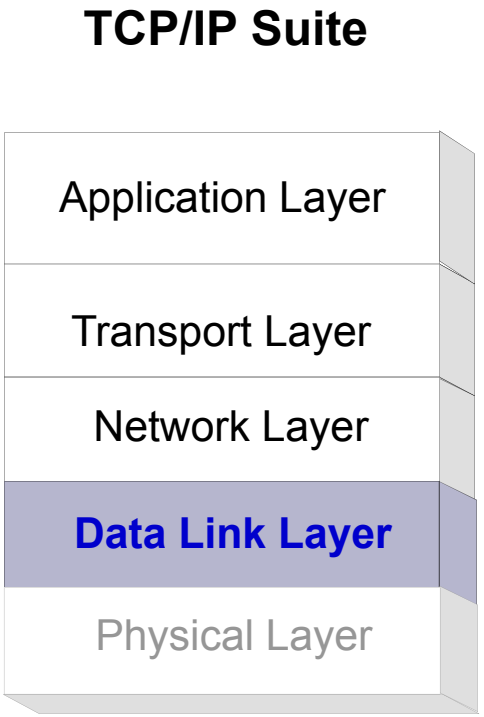

Types of Networks

Point-to-point network

- Each link connects two end points: hosts or any network devices/elements
- Usually for long distance connections
- Examples: DSL (digital Subscriber Loop), POS (Packet over SONET/SDH), GbE (Gigabit Ethernet)

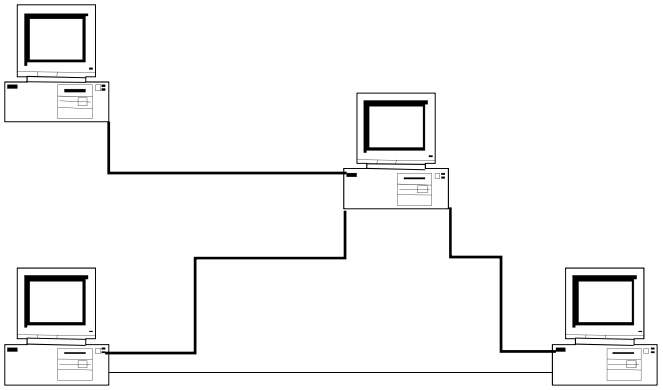

Point-to-Point Network

Broadcast network

- A number of stations share a common transmission medium
- Usually for local networks
- Examples: CSMA/CD Ethernet, WLAN (Wireless Local Area Network), a.k.a. Wi-Fi

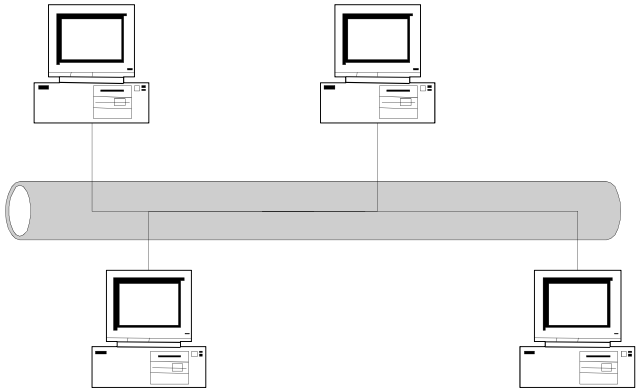

Broadcast Network

Multiple Access

- Network topology
  - Point-to-point → N(N-1)/2 links to connect N nodes
  - Broadcast → the shared medium forms a single domain
- Medium access control (MAC) protocols
  - Rules to share a medium
  - Carrier Sense Multiple Access/Collision Detection (CSMA/CD)
  - Carrier Sense Multiple Access/Collision Avoidance (CSMA/CA)

Point-to-Point Protocol
– point-to-point network example

- The Point-to-Point Protocol (PPP) is a data link protocol
- The main purpose of PPP is encapsulation and transmission of IP datagrams, or other network layer protocol data, over a serial link.
- Currently, PPP is used by most dial-up Internet access, Digital Subscriber Loop (DSL), and cable broadband services.
- DSL example

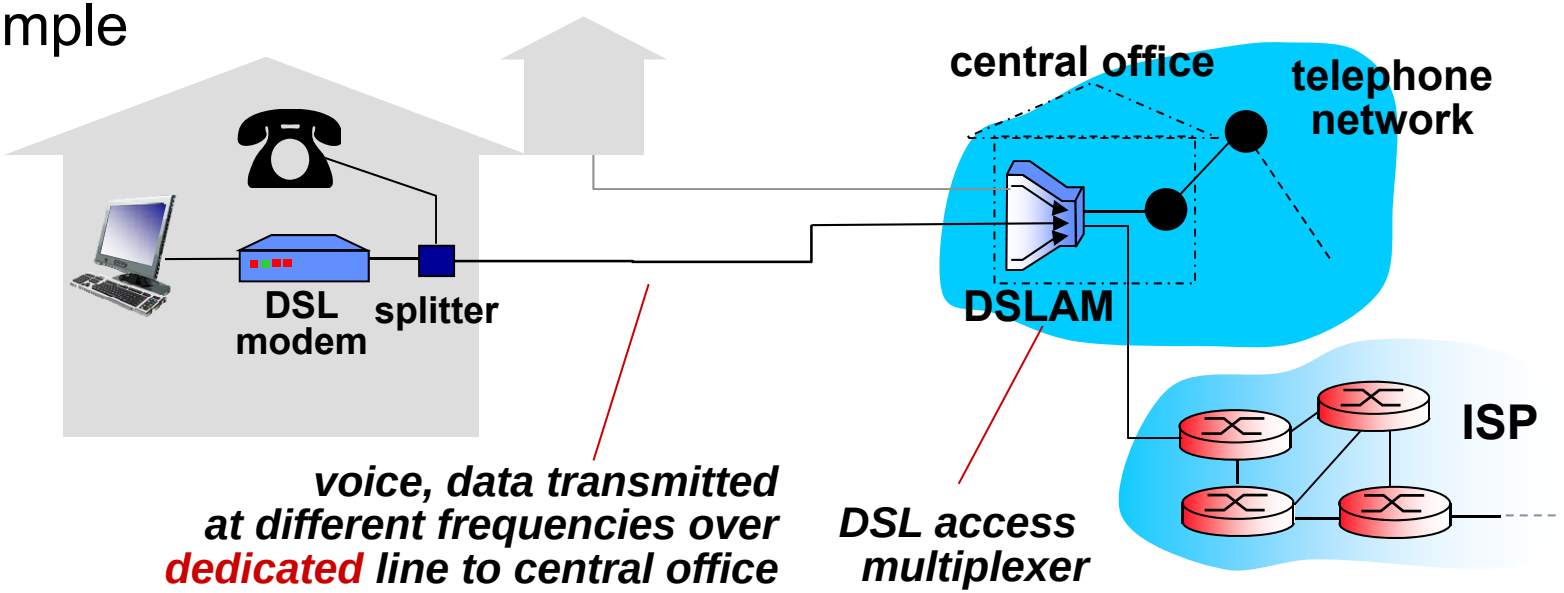

PPP Encapsulation

PPP frame format

- Flag: mark the beginning and ending of a frame
- Protocol: used to multiplex different protocol data
- No addressing, only two end hosts.

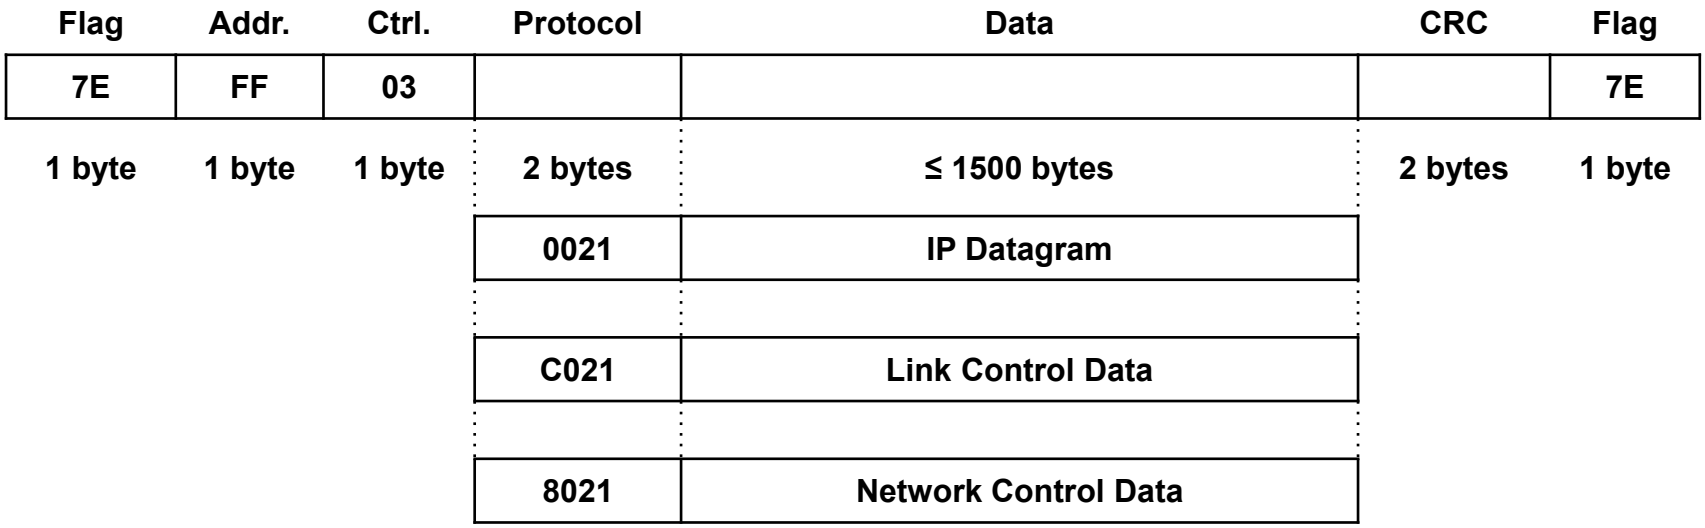

Point-to-Point Protocol (RFC 1661)

PPP consists of two types of control protocols:

- Link Control Protocol (LCP)
  - Responsible for agreeing on PPP encapsulation options, packet size limits, and detecting common mis-configuration errors over the data link
  - Optional features to provide peer authentication, detect link status
- Network Control Protocol (NCP)
  - PPP supports a family of NCPs and treat each network protocol like an interface
  - IP Control Protocol (IPCP, RFC 1332), used for configure the link to transmit IP datagrams

Local Area Networks

- Local Area Networks (LANs) typically connect computers within a building or a campus.
- Many LANs are broadcast networks.
- Bus and Ring are two typical LAN topologies used in early days
- The protocol that determines who can transmit on a broadcast channel is called Medium Access Control (MAC) protocol.

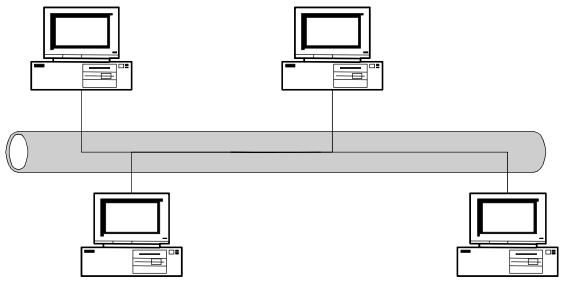

Bus LAN

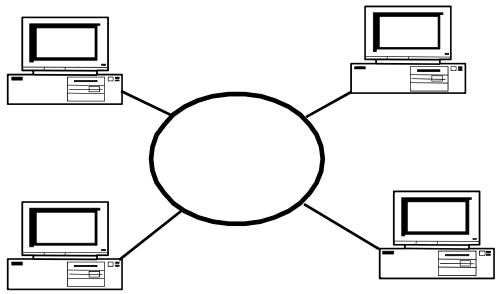

Ring LAN

MAC and LLC

- In any broadcast network, the stations must ensure that only one station transmits at a time on the shared communication channel.
- The protocol that determines who can transmit on a broadcast channel is called Medium Access Control (MAC) protocol.
- The MAC protocol is implemented in the MAC sublayer which is the lower sublayer of the data link layer.
- The higher portion of the data link layer is often called Logical Link Control (LLC).

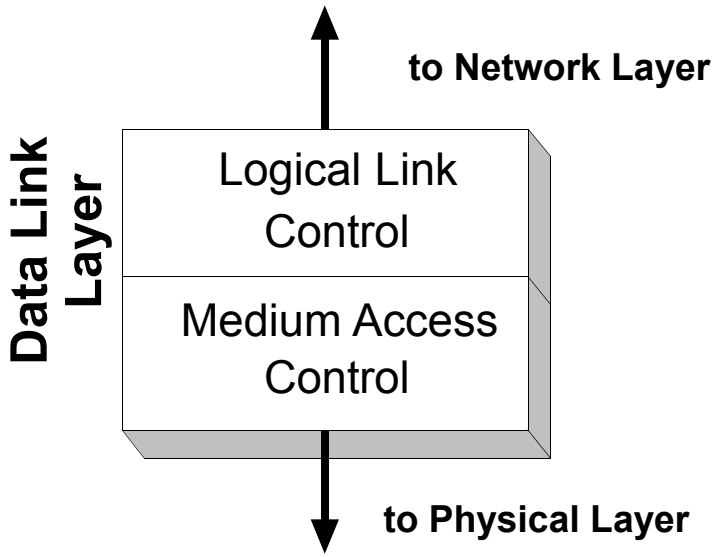

Logical Link Control

- LLC can provide different services to the network layer:
  - unacknowledged connectionless service
  - acknowledged connectionless service
  - connection-oriented service
- Framing
- Error control
- Addressing

Media Access Control

MAC algorithms are used to resolve collisions and share the medium in a broadcast network.

Examples of MAC:

- ALOHA
- Carrier Sense Multiple Access/Collision Detection (CSMA/CD)
- Carrier Sense Multiple Access/Collision Avoidance (CSMA/CA)

Ethernet

- Originally developed at Xerox PARC at 2.94 Mbps data rate over coaxial cable based on CSMA/CD, 1972 – 1977
- An industry standard since 1982 (DIX Ethernet), or
- Based on CSMA/CD adopted by IEEE 802 committee in 1985
  - Interoperates with 802.2 (LLC) as higher layer
  - Uses different encapsulation header to carry data payload

CSMA/CD Flow Diagram

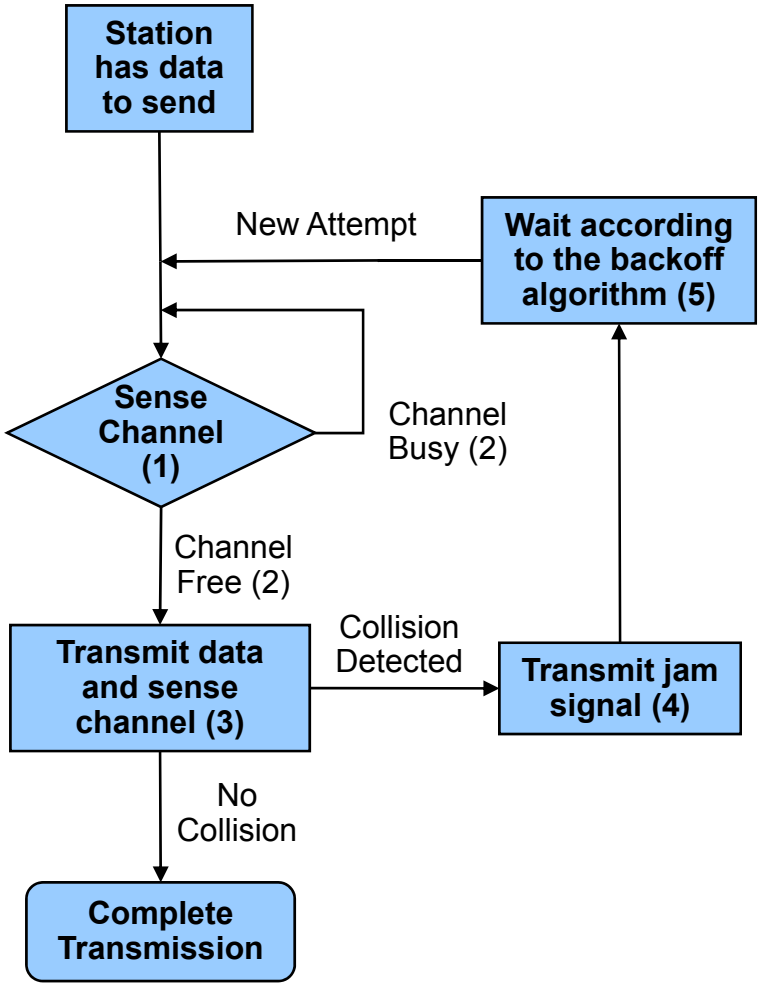

1. Each station listens before it transmits.
2. If the channel is busy, it waits until the channel goes idle, and then it transmits.
3. If the channel is idle it transmits immediately and continue sensing for ***2a*** seconds.
4. If collision is detected, transmit a brief jamming signal then cease transmission.
5. Wait for a random time, and retransmit. The random time is determined by exponential backoff algorithm.

CSMA/CD (collision detection)

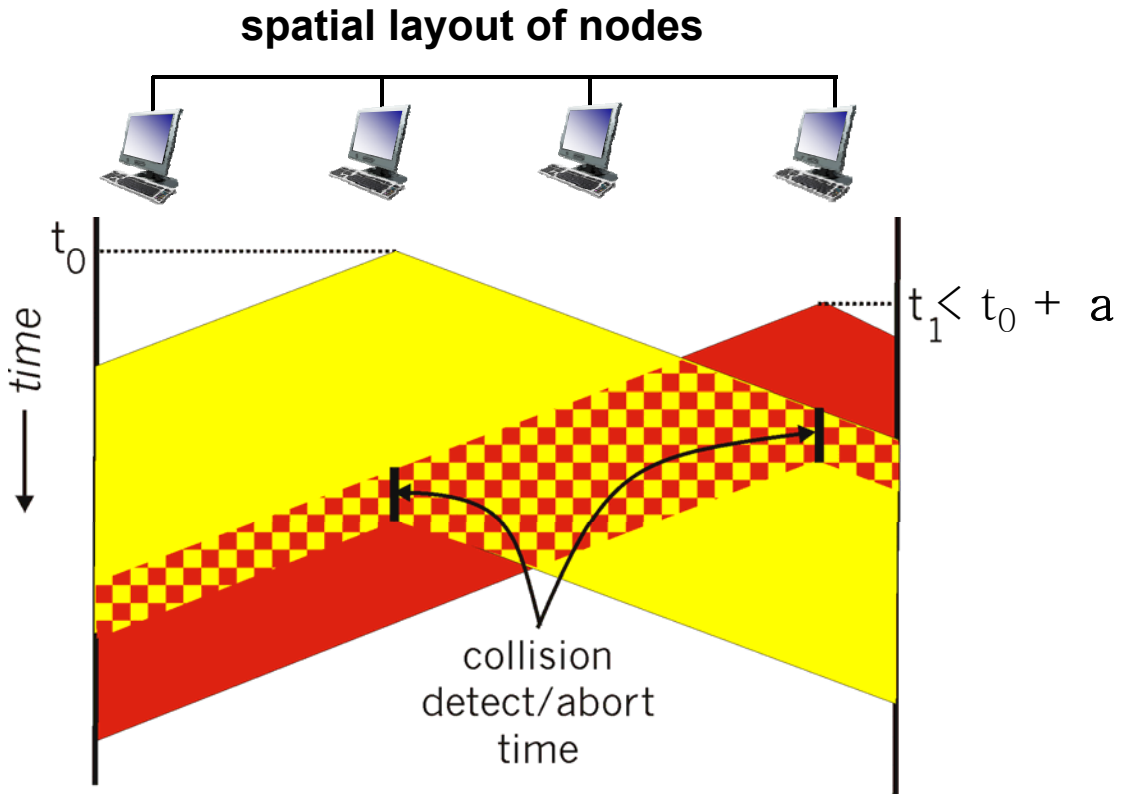

Collisions in Ethernet

- The collision resolution process of Ethernet requires that a collision is detected while a station is still transmitting.
- Assume the maximum propagation delay on the bus is *a*.
- Restrictions: Each frame should be at least twice as long as the time to detect a collision (*2a*).

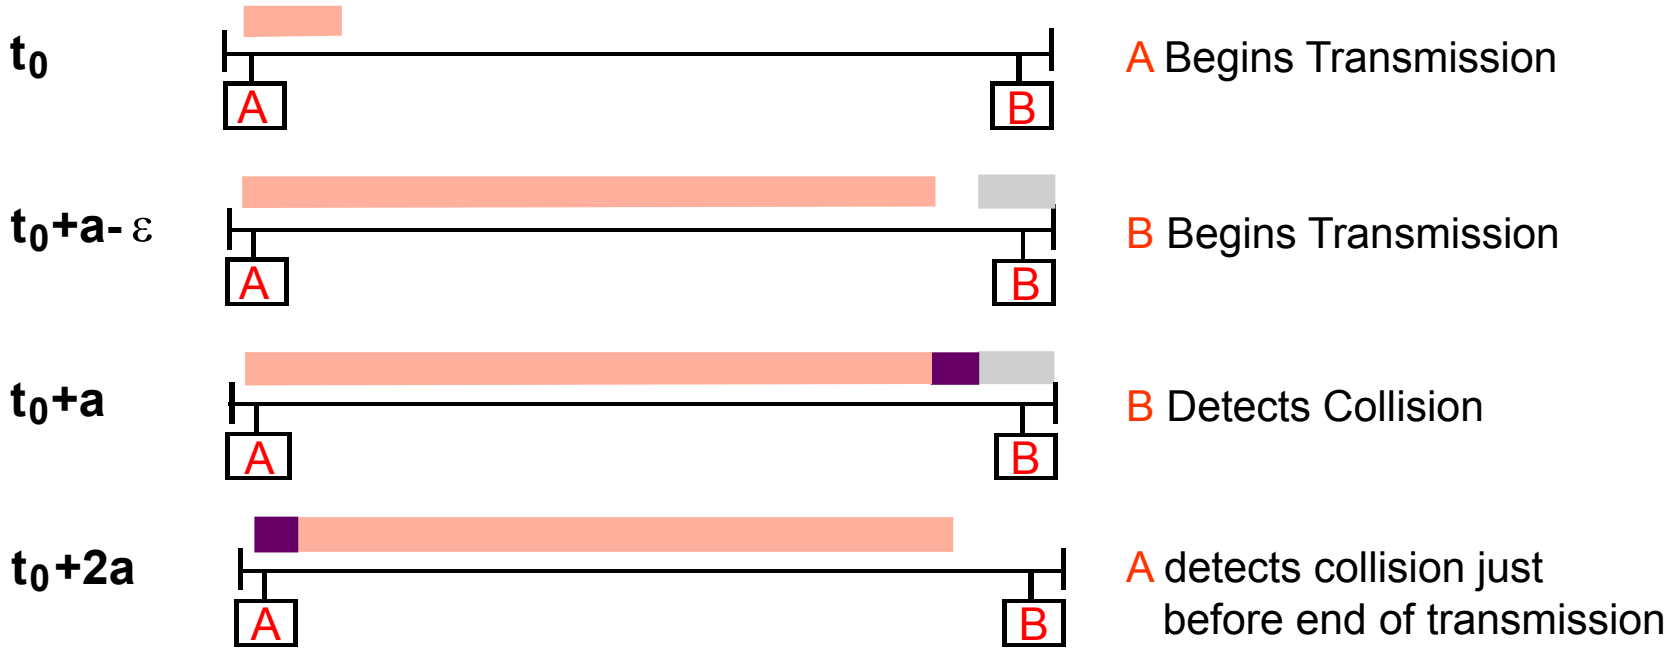

Exponential Backoff Algorithm

- If a station is involved in a collision, it waits a random amount of time before attempting a retransmission.
- The random time is determined by the following algorithm:
  - Set "slot time" to *2a*.
  - After first collision wait *0* or *1* time unit.
  - After the i^{th} collision, wait a random number between *0* and 2^i*-1* time slots.
  - Do not increase random number range if *i>9*.
  - Give up after *16* collisions.

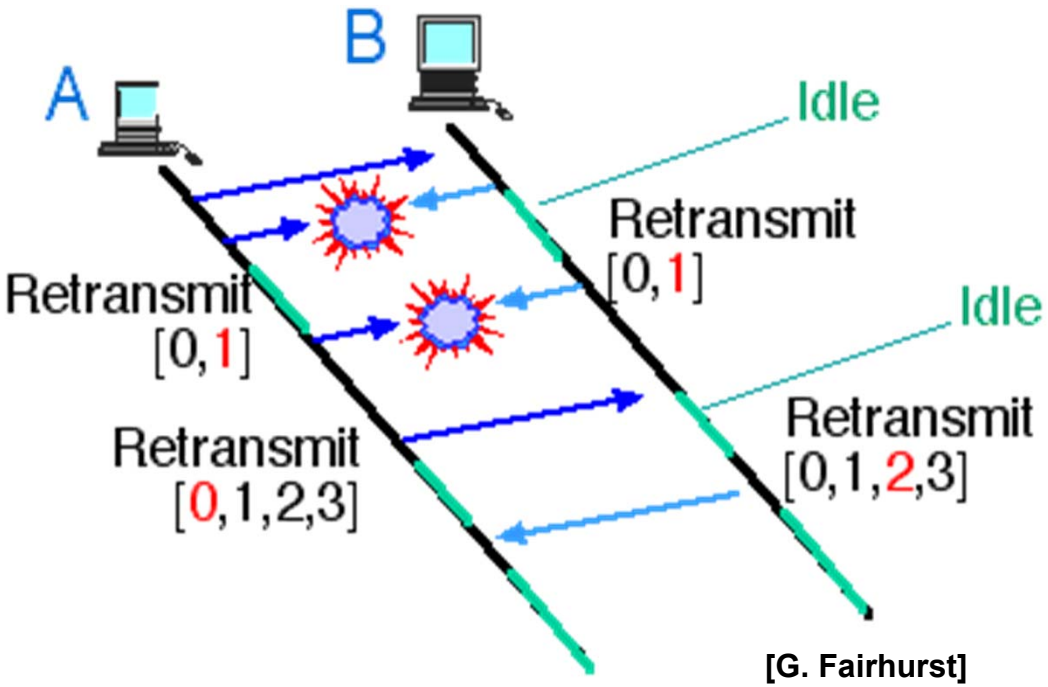

[G. Fairhurst]

Ethernet Switches

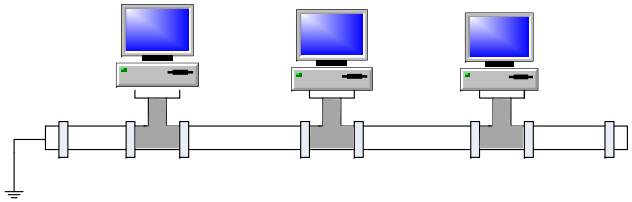

In an Ethernet LAN, hosts can be

- Attached to a common cable, or
- Connected by *Ethernet switches*.

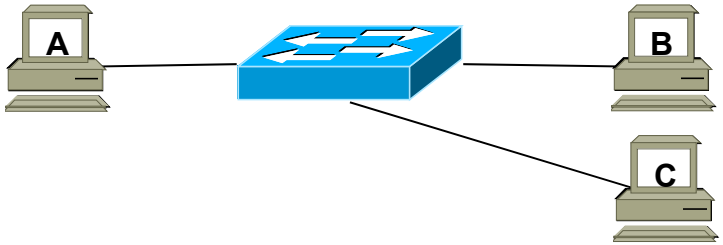

Ethernet switches are MAC layer devices that switch frames between ports connected to different LAN segments.

- Offer guaranteed bandwidth for segments.
- Separate a LAN into collision domains.

Ethernet Encapsulation (RFC 894)

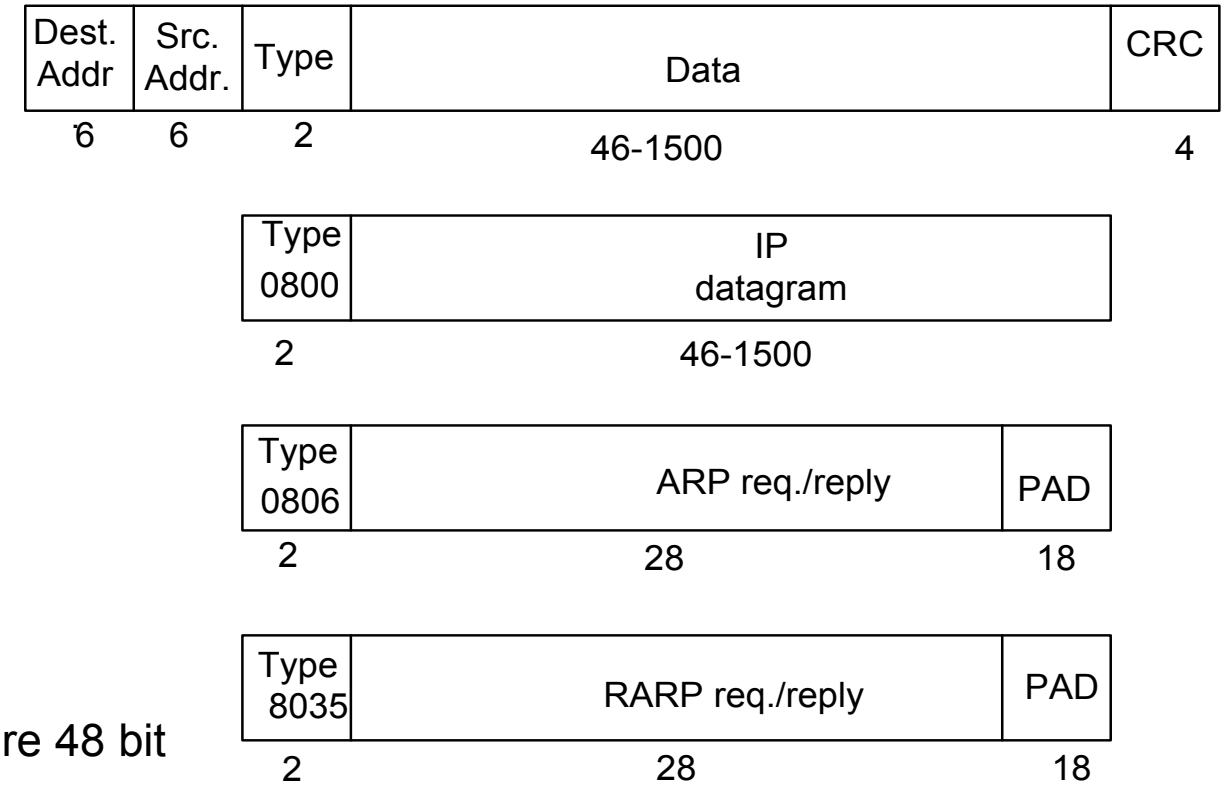

- **Dest. Addr., Src. Addr.:** MAC addresses are 48 bit
- **Type:** Identifies the content of the data field (must ≥ Ox0600)
- **CRC:** cyclic redundancy check

IEEE 802.2/802.3 Encapsulation

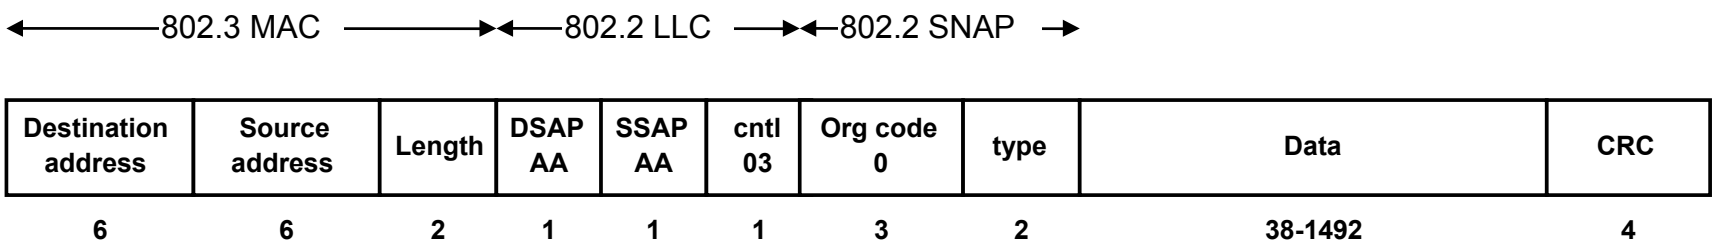

- Destination address, Source address:
  MAC addresses are 48 bit (displayed as 12 hexadecimal characters)
- Length: frame length in number of bytes (<Ox0600, 1,500 bytes → Ox05dc)
- DSAP, SSAP: always set to 0xaa
- Ctrl: set t o 3
- Org code: set to 0
- Type field: identifies the content of the data field
- CRC: cyclic redundancy check

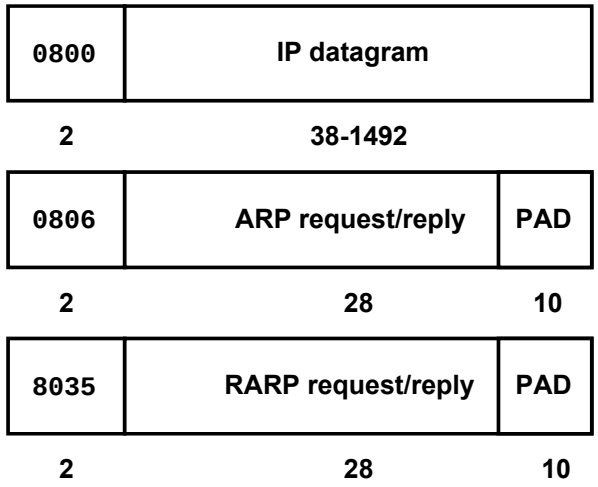

Total frame size: 64 bytes to 1518 bytes

Overhead: 38 bytes including 12 bytes Inter Frame Gap (IFG)

The Address Resolution Protocol

- IP addresses are not recognizable in the interface layer where physical addresses (or MAC addresses) are used.
- Different kinds of physical networks use different addressing schemes.
- Address Resolution Protocol (ARP): maps an IP address to a MAC address per RFC 826.
- Reverse Address Resolution Protocol (RARP): maps a MAC address to an IP address per RFC 903.

ARP Process

- When a source host wants to send an IP packet to a destination, it first *broadcasts* an *ARP request* asking for the MAC address corresponding to a target IP address.
- A target device will return an *ARP reply* with its MAC address.

ARP Packet Format

Hardware Type	Protocol Type	Hardware Size	Protocol Size	Operation Field	Sender Eth. Addr.	Sender IP Addr.	Target Eth. Addr.	Target IP Addr.
2	2	1	1	2	6	4	6	4 bytes

- 28 bytes long.
- An *ARP request* or *ARP reply* is encapsulated in an Ethernet frame.
  - The protocol type in Ethernet frame is set to 0x0806 for ARP messages.
- Hardware Type - Specifies a hardware interface type for which the sender requires a response, i.e. Ethernet (1) in our case
- Protocol type - Specifies the type of high-level protocol address the sender has supplied, ie. IP (Ox0800) in our case
- Hlen - Hardware address length.
- Plen - Protocol address length.
- Operation field specifies ARP request (1), ARP reply (2), RARP request (3), or RARP reply (4).

ARP Request

Hardware Type	Protocol Type	Hardware Size	Protocol Size	Operation Field	Sender Eth. Addr.	Sender IP Addr.	Target Eth. Addr.	Target IP Addr.
2	2	1	1	2	6	4	6	4 bytes

- Ethernet destination: ff:ff:ff:ff:ff:ff (broadcast address)
- Target Ethernet Address: not set.

ARP Reply

Hardware Type	Protocol Type	Hardware Size	Protocol Size	Operation Field	Sender Eth. Addr.	Sender IP Addr.	Target Eth. Addr.	Target IP Addr.
2	2	1	1	2	6	4	6	4 bytes

The ARP reply is sent by the node whose IP address matches the *target IP address* in the ARP request.

- It fills its MAC address into the *target Ethernet address* field of the ARP request.
- It then swaps the two sender addresses (Ethernet and IP addresses) with the two target addresses, sets the op field to 2.
- The ARP reply is sent back to the source host only.

All other nodes receiving the broadcast ARP ignore the request, since their IP addresses do not match the target IP address.

ARP Cache

- Sending an ARP request/reply for each IP datagram is inefficient.
- Each host maintains an ARP cache containing the recent resolved IP addresses.
- A source host first checks its ARP cache for the destination MAC address,
  - If an entry is found, sends out the IP packet within an Ethernet frame.
  - Otherwise, sends out an ARP request.

Manipulating the ARP Table

- Elements of an entry in the ARP table:
  - An IP address
  - A MAC address
  - Flags
- A normal entry expires after 20 minutes after it is created or the last time it is referred.
- Manipulate ARP table by the arp command:
  - arp –a: Displays all entries in the ARP table.
  - arp –d: Deletes an entry in the ARP table.
  - arp –s: Inserts an entry into the ARP table.

Proxy ARP (RFC 1027)

- Hide the two physical networks from each other.
- Use a proxy-ARP-enabled router to answer ARP requests targeted for a host …

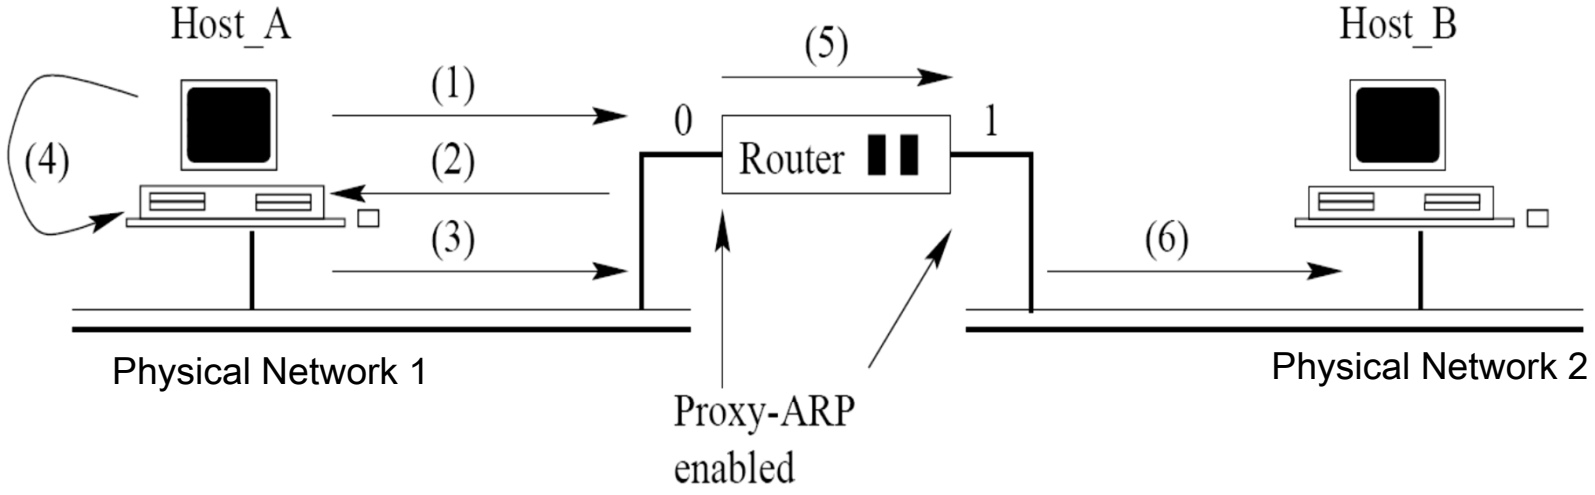

(1): Host_A sends ARP request for Host_B's MAC

(2): Router Port 0 replies for Host_B

(3): Host_A sends the frame to Router Port 0

(4): Host_A inserts a new entry in its ARP cache: {(Host_B's IP) at (Router Port 0's MAC)}

(5): Router forwards the frame to port 1

(6): Router port 1 sends the frame to Host_B

Note

- Two networks are logically in the same subnet (at least from Host A's point of view …)
- Often used by a router connecting to a host with a serial link, ex. PPP to Host B

Gratuitous ARP

- Occurs when a host sends an ARP request resolving its own IP address.
- Usually happens when the interface is configured at bootstrap time.
- The interface uses gratuitous ARP to determine if there are other hosts using the same IP address.
- The sender's IP and MAC address are broadcast, and other hosts will insert this mapping into their ARP tables.

Loopback Interface

- Most TCP implementations have a loopback interface with IP address 127.0.0.1 and named as localhost.
- The localhost behaves as a separate data link interface.
- A packet that is sent to the loopback interface moves down the protocol stack and is returned back by the driver software for the localhost “device”.
- Used for debugging.
- Packets sent to loopback interface will not appear on network.

Network Interface Operations

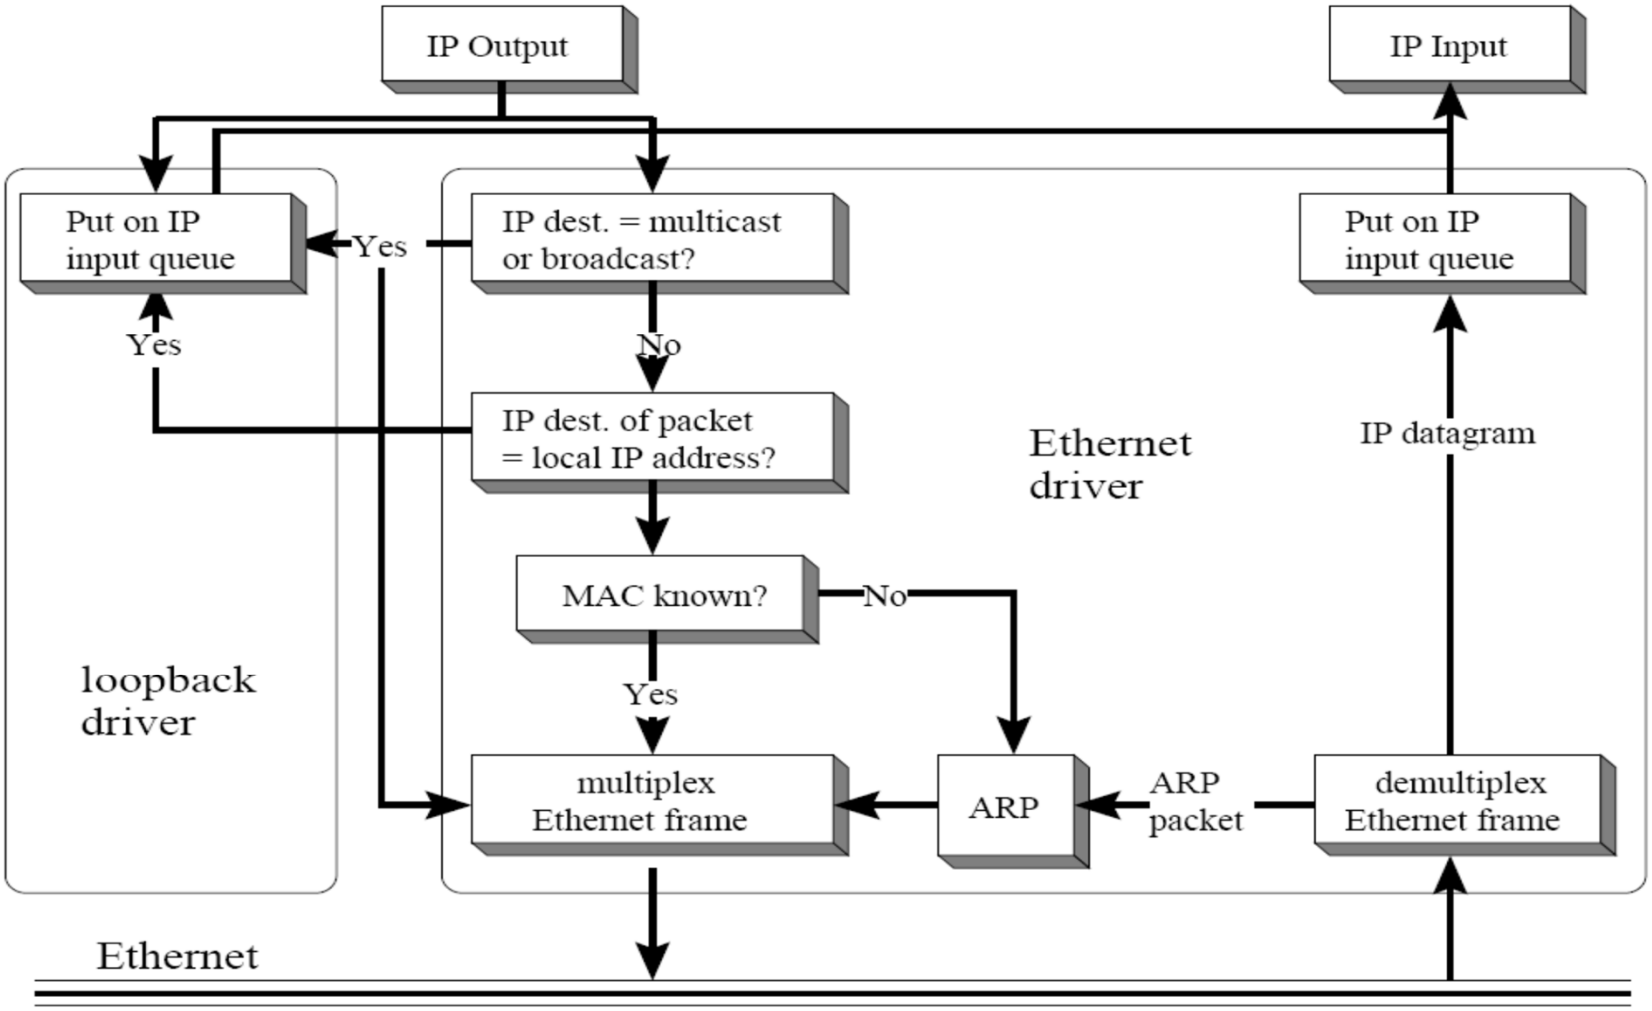

Functional Diagram of an Ethernet Interface Card

Maximum Transmission Unit

- There is a limit on the data packet size of each data link layer protocol.
- This limit is called Maximum Transmission Unit (MTU).
- MTUs for various data link layers:
  - Ethernet, PPP: 1500 bytes
  - FDDI: 4352 bytes
  - PPP (low delay): 296 bytes
- MTU does not count its own header and trailer bytes of the data link protocol. e.g. Ethernet's MTU is 1500 bytes.

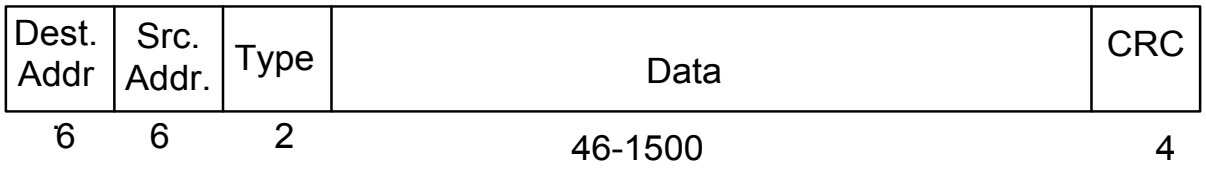

Internet Control Message Protocol

- The Internet Control Message Protocol (ICMP) is the protocol used for error and control messages in Internet.
- ICMP provides an error reporting mechanism of routers to the sources.
- All ICMP packets are encapsulated as IP datagrams (IP protocol type 1)
- The packet format is simple.

Example below shows destination unreachable error message below:

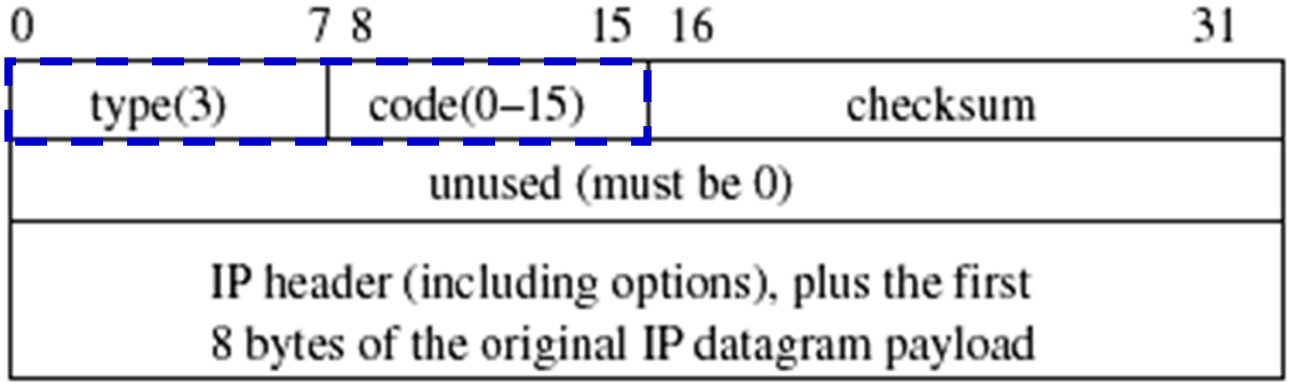

Types of ICMP Packets

Many ICMP packet types exist, each with its own format.

Type	Code	Description	–
0	0	echo reply	query
3	0–15	destination unreachable	error
5	0–3	redirect	error
8	0	echo request	query
9	0	router advertisement	query
10	0	router solicitation	query
11	0–1	time exceeded	error

ICMP Message Types

- ICMP messages are either query messages or error messages.
- ICMP query messages:
  - Echo request / Echo reply
  - Router advertisement / Router solicitation
  - Timestamp request / Timestamp reply
  - Address mask request / Address mask reply
- ICMP error messages:
  - Host unreachable
  - Source quench
  - Time exceeded
  - Parameter problem

ICMP Error Messages

- Each ICMP error message contains the header and at least the first 8 bytes of the IP datagram **payload** that triggered the error message.
- To prevent that too many ICMP messages, ICMP error messages are NOT sent
  - for multiple fragments of the same IP datagrams
  - in response to an error message
  - in response to a broadcast packet

ICMP Type 3 Error Message Code

There are 16 different ICMP error messages ('codes') of type "Destination Unreachable" (Type = 3)

Code	Message Type	Code	Message Type
0	Network unreachable	8	Source host isolated
1	Host unreachable	9	Destination network administratively prohibited
2	Protocol unreachable	10	Destination host administratively prohibited
3	Port unreachable	11	Network unreachable for TOS
4	Fragmentation needed but bit not set	12	Host unreachable for TOS
5	Source route failed	13	Communication administratively prohibited by filtering
6	Destination network unknown	14	Host precedence violation
7	Destination node unknown	15	Precedence cutoff in effect

ICMP Port Unreachable

If, in the destination host, the IP module cannot deliver the datagram because the indicated protocol module or process port is not active, the destination host may send a port unreachable message to the source host.

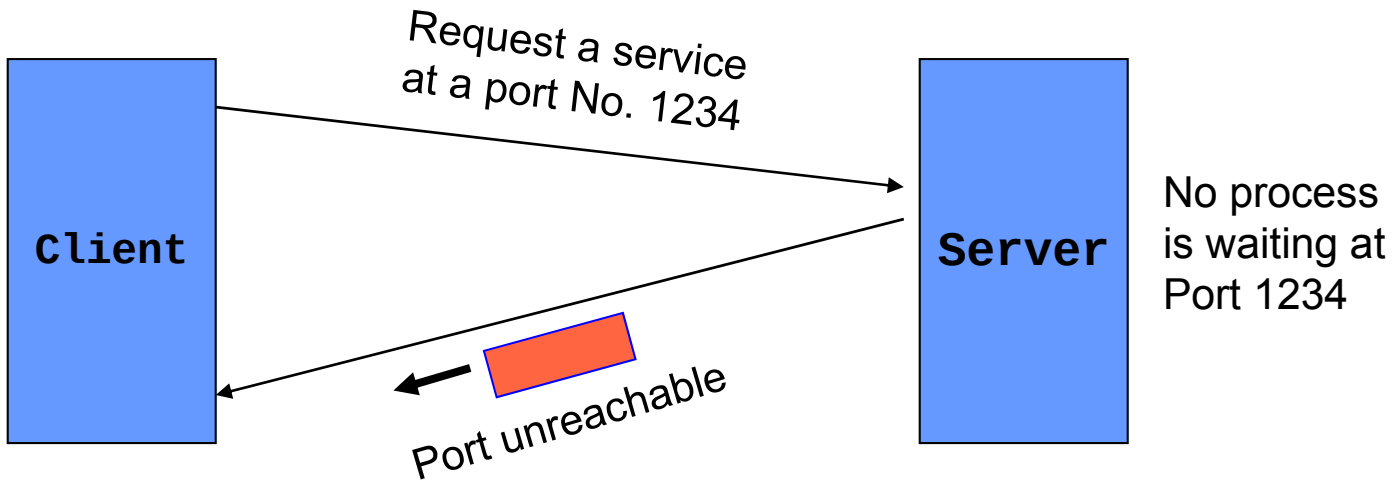

ICMP Port Unreachable

Format of the Port Unreachable Message

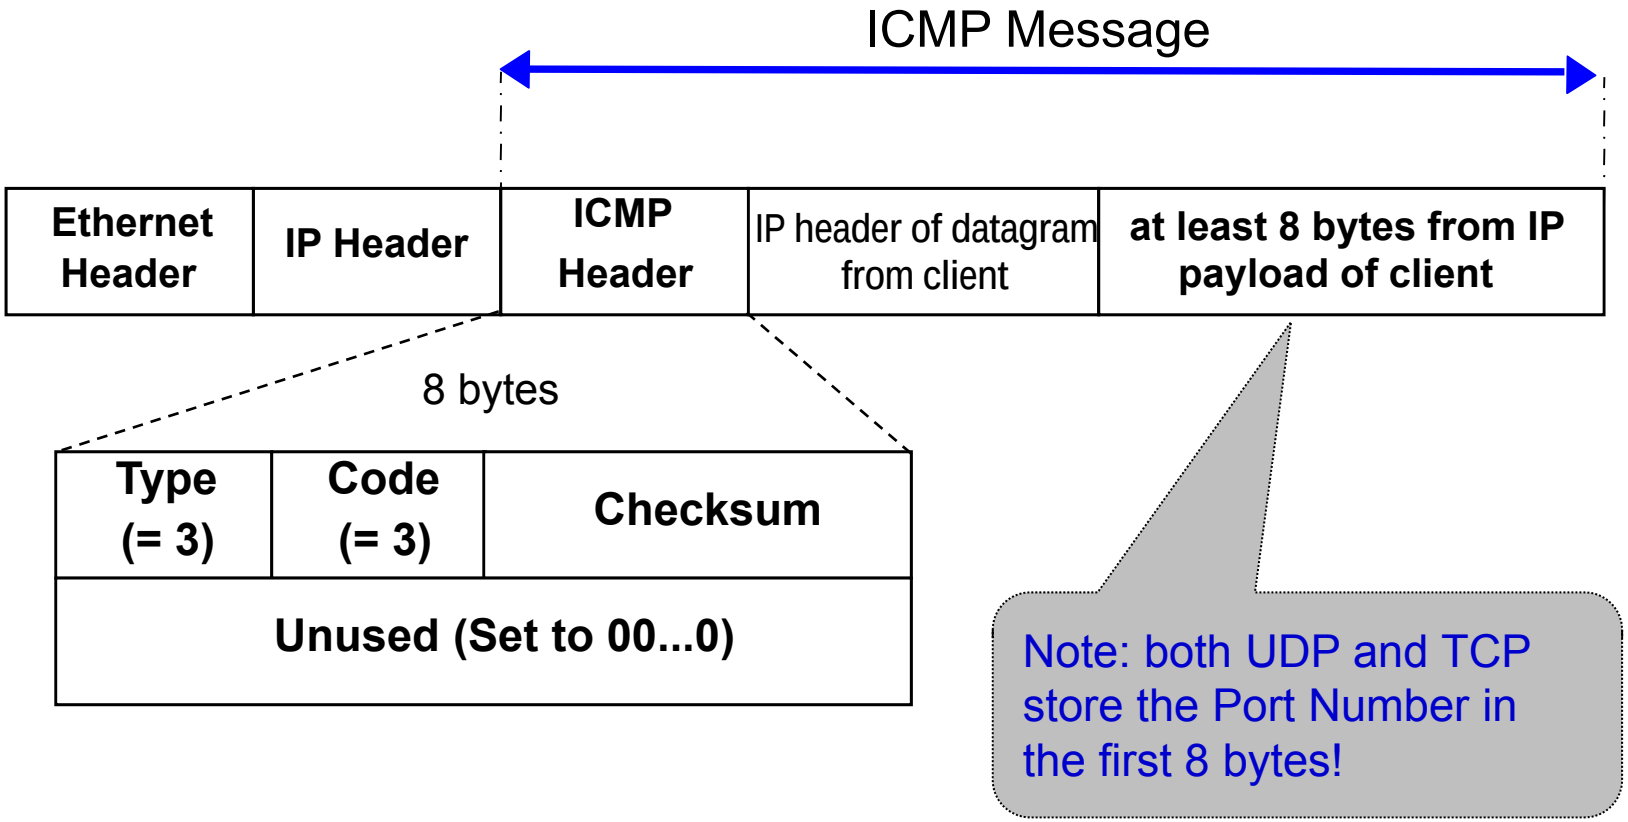

Packet InterNet Gopher (PING)

- PING is a program that utilizes the ICMP echo request and echo reply messages.
- PING is used to verify if a certain host is up and running. It is used extensively for fault isolation in IP networks.
- PING can be used with a wide variety of options, e.g.
  - *-R* (Record route): includes the RECORD_ROUTE option in the ECHO_REQUEST packet and displays the route buffer on returned packets.
  - *-s packetsize*: specifies the number of data bytes to be sent (default is 56) (in newer implementations, -s is used to continuously generate queries)

Echo Request and Reply

- Ping's are handled directly by the kernel.
- Each Ping is translated into an ICMP Echo Request.
- The Ping'ed host responds with an ICMP Echo Reply.

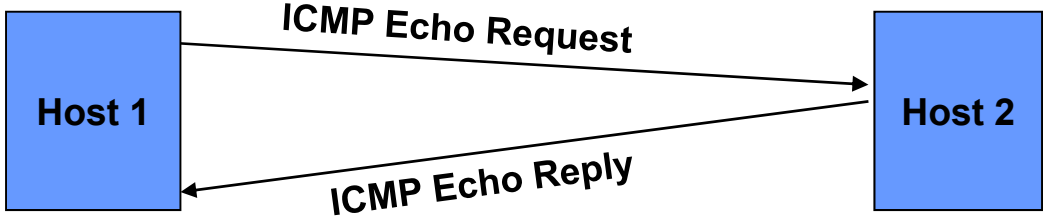

- Message format
  - Identifier is set to process ID of querying process.
  - Sequence number is incremented for each new echo request.

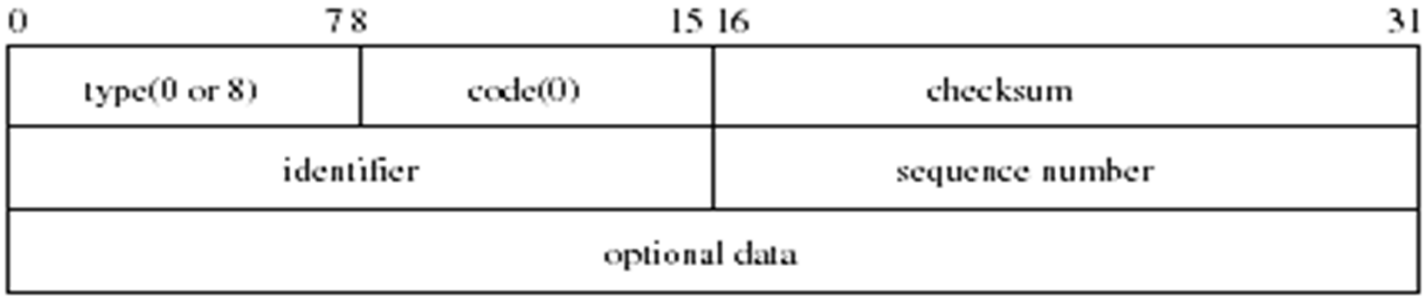

Sock Traffic Generator

- Sock is a test program.
  - Can be run as a client or as a server
  - Use UDP or TCP.

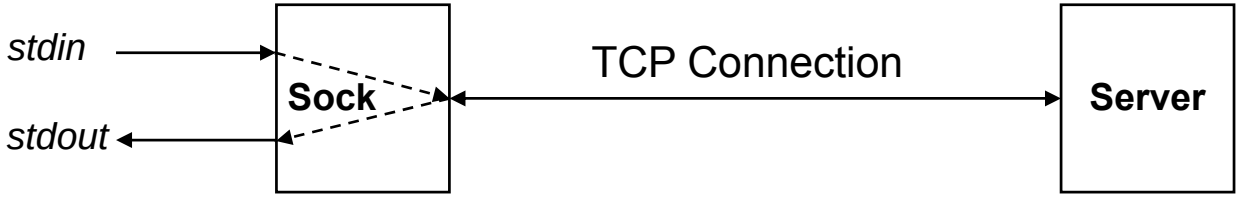

- Sock operates in one of the following four modes:
  - Interactive client
    - > Copy data from stdin, send to server, receive data back from server, copy on screen (stdout)
  - Interactive server
    - > Wait for a request from client
  - Source client
    - > Send packets to a specific server
  - Sink server
    - > Receive packets from a client and discard the data