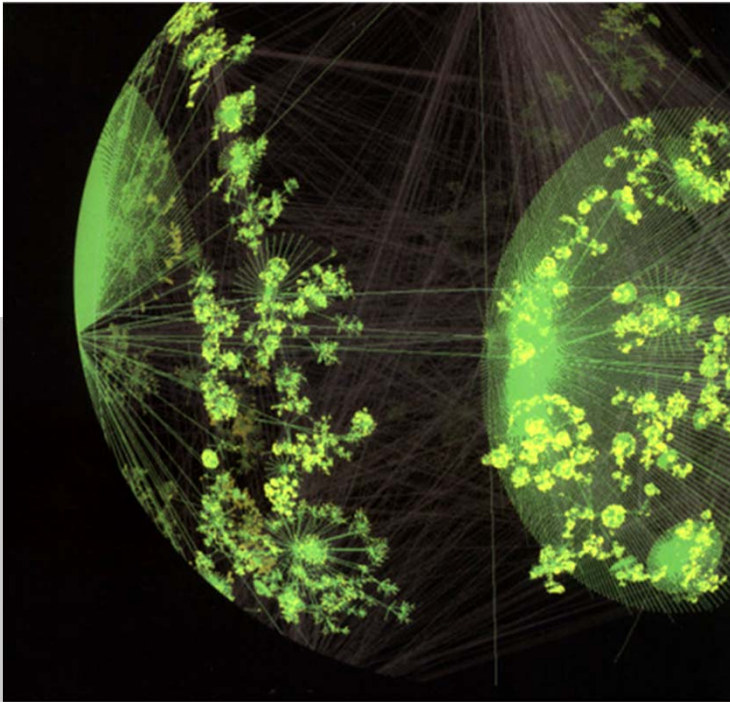


Chapter 5

UDP and Its Applications

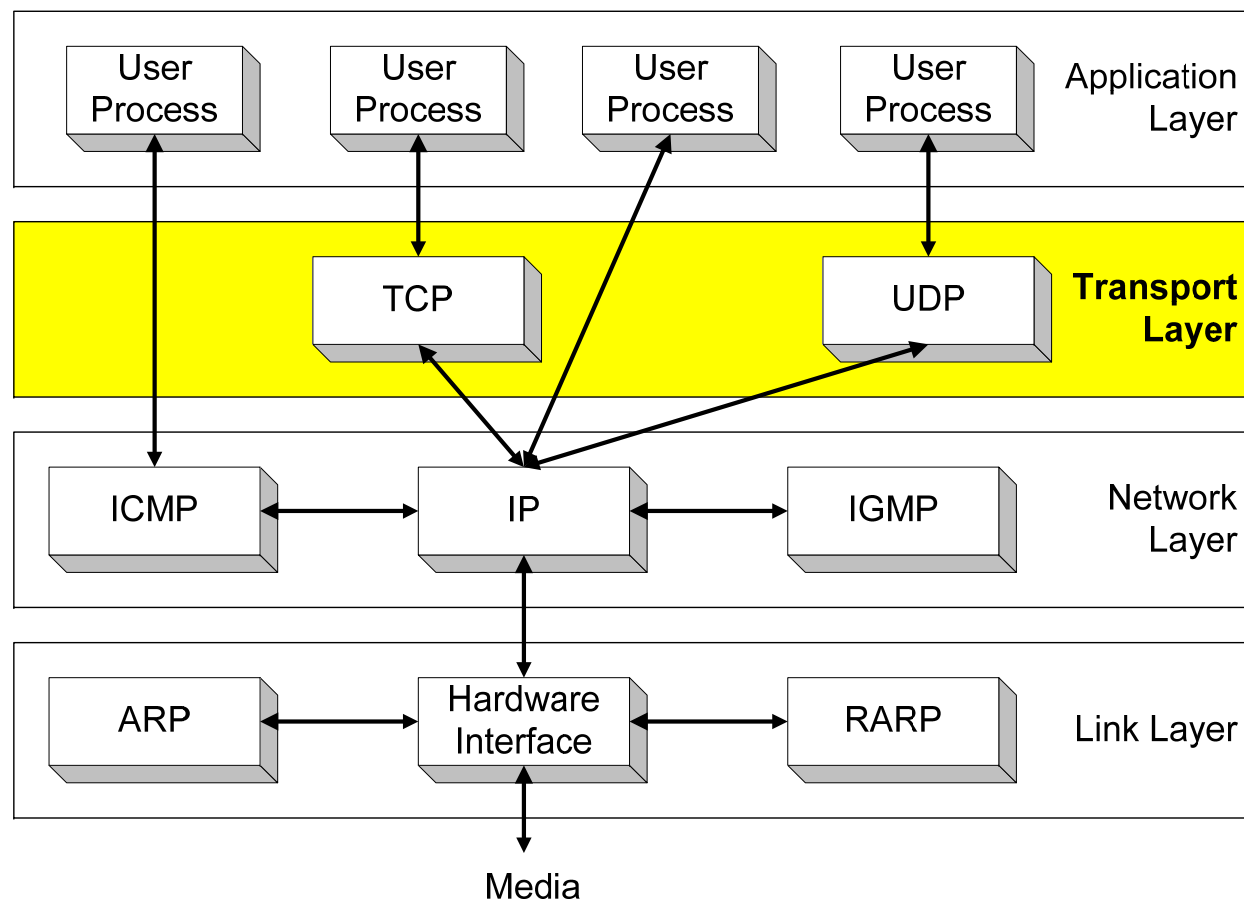


TCP/IP Essentials
A Lab-Based Approach

Spring 2017


Orientation

The transport layer



Transport Layer Protocols

- Transport layer protocols are end-to-end protocols
- They are only implemented at the hosts



UDP and TCP



The Internet supports two transport protocols:

UDP – User Datagram Protocol


- Datagram oriented
- Unreliable, connectionless
- Simple
- Unicast and multicast
- Commonly used for network control signaling services
 - Network management (SNMP), routing (RIP), naming (DNS), etc.
- Useful for increasing number of applications, e.g., multimedia applications

TCP – Transmission Control Protocol

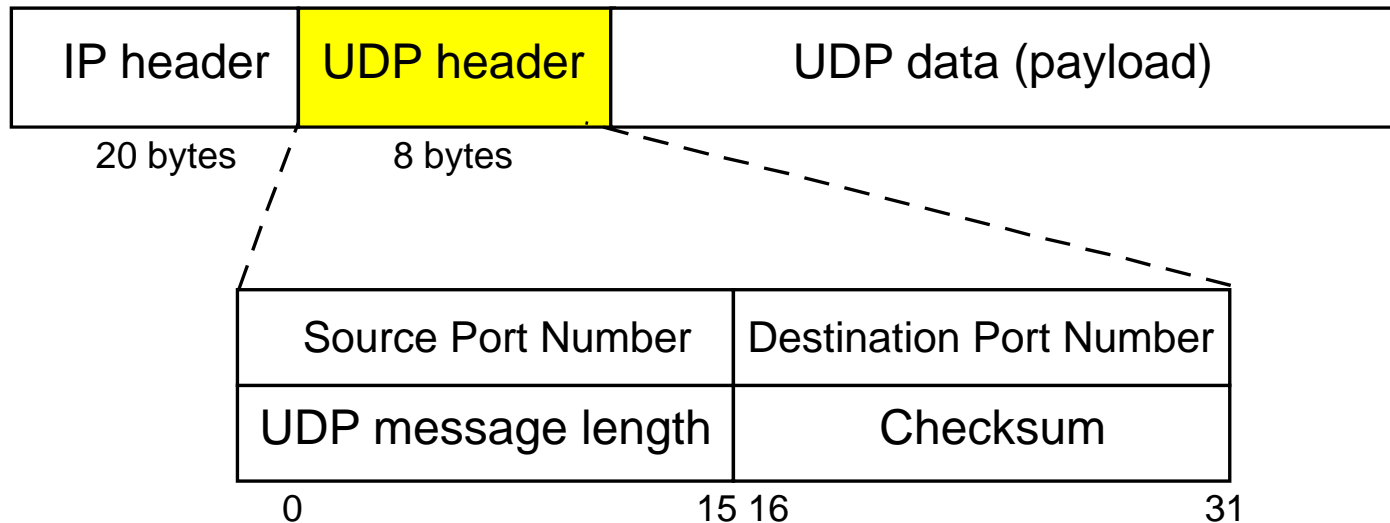
- Stream oriented
- Reliable, connection-oriented
- Complex
- Unicast only
- Currently used by most Internet applications:
 - Web (HTTP), email (SMTP), file transfer (FTP), terminal (telnet), etc.

UDP - User Datagram Protocol

- UDP supports unreliable transmissions of datagrams
- UDP merely extends the host-to-host delivery service of IP datagram to an application-to-application service
- The only thing that UDP adds is multiplexing and demultiplexing



UDP Format



- **Port Numbers** identify sending and receiving applications (processes). The maximum value for a port number is $2^{16}-1= 65,535$
- **Message Length** is between 8 bytes (i.e., data field can be empty) and 65,535 bytes (length of UDP header and data in bytes)
- **Checksum** is for UDP header and UDP data

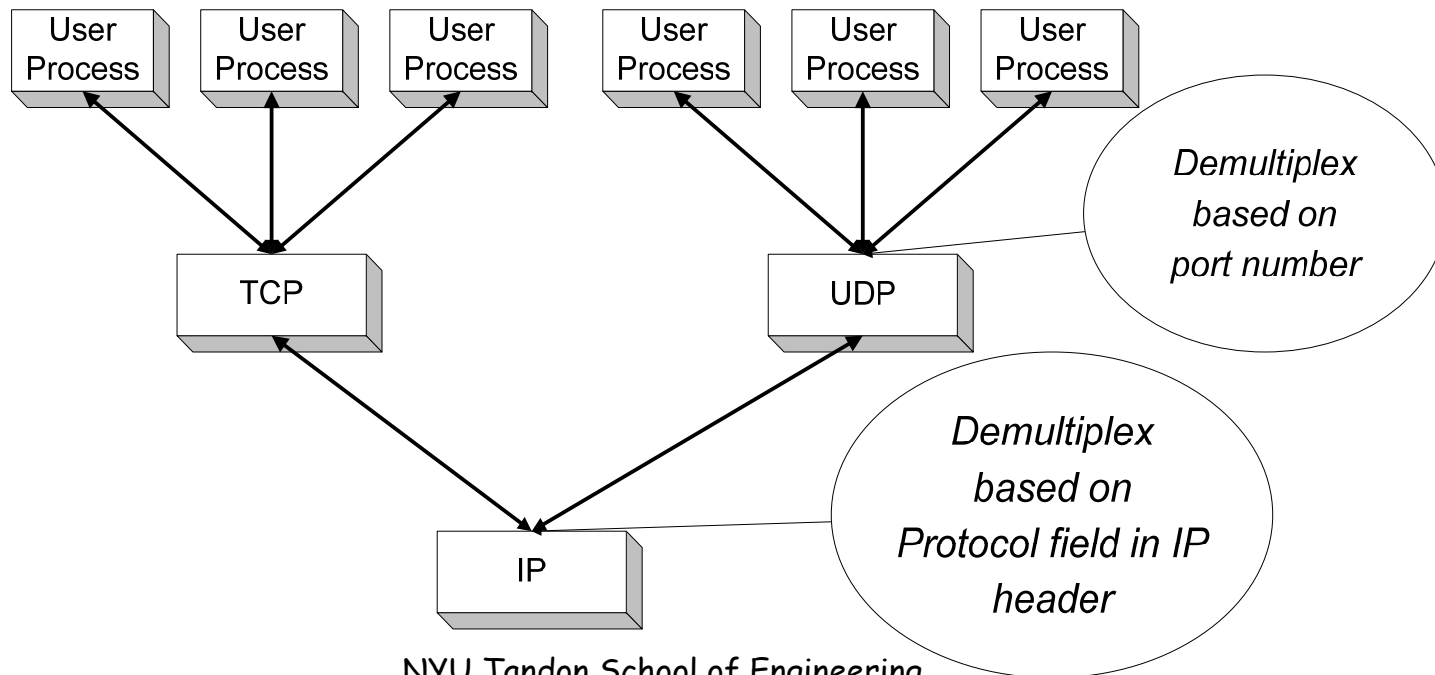
UDP Checksum

- Optional
 - set all 0's if not calculated
 - A calculated checksum can never be all 0's.
- Computed using the UDP header, UDP data and a pseudo-header as below.
- All fields of pseudo-header are available in UDP layer

32-bit Source IP Address		
32-bit Destination IP Address		
0x00	8-bit Protocol (0x17)	16-bit UDP Length (bytes)

Port Numbers

- UDP (and TCP) use port numbers to identify applications
- A globally unique flow of host application can be identified by a 5-tuple **<Src. IP, Dst IP, Src. Port, Dst. Port, Protocol No.>**
- There are 65,535 UDP ports available per host
 - Dynamic/private , used by clients, randomly picked, >49,151 (per IANA)
 - Registered, used by ordinary user processes, 1024 – 49,151
 - Well-known, used by servers, fixed, 1~1023



Ephemeral Port Range

Ephemeral Port: short-lived port used as transport protocol port automatically allocated from a predefined range by the TCP/IP stack software.

- The IANA suggests 49,152 to 65,535 as "dynamic and/or private ports."
- The BSD uses ports 1,024 through 4,999 as ephemeral ports
- Many Linux kernels use 32,768 to 61,000 specified in */proc/sys/net/ipv4/ip_local_port_range*
- MS Windows OS' through Server 2003 use the range 1,025 to 5,000 as ephemeral ports; use the IANA range since Windows Vista and Server 2008
- FreeBSD uses the IANA port range since release 4.6.

About Port 0



- Port 0 is a reserved port by IANA
- Many OS' allows a source port of 0 from a high layer application for connecting to a remote host. The OS automatically reassigns an ephemeral port when encapsulating the application data to a transport segment
- No traffic should flow over Internet using port 0 although different OS' have different ways of handling traffic using port 0
- The external behavior of handling port 0 in different OS's can be “fingerprinted” with a set of tests to send TCP or UDP packet from source port 0 to different destination ports
- It is highly recommend that one should block any traffic using this port at your firewall so no program should be listening on port 0 and no program should connect from port 0


Maximum Transmission Unit (MTU)

- The frame size limit of the data link protocol specifies a limit on the size of the IP datagram that can be encapsulated by the protocol.
- This limit is called **Maximum Transmission Unit** (MTU)
- MTUs various for different data link protocols:

Ethernet: 1500	FDDI: 4352
802.2/802.3: 1492	ATM AAL5: 9180
802.5: 4464	PPP: 296 (low delay)
- What if the size of an IP datagram exceeds the MTU?
 - IP datagram is fragmented into smaller units.
- What if the route contains networks with different MTUs?
 - The smallest MTU of any data link is used as the **Path MTU**.

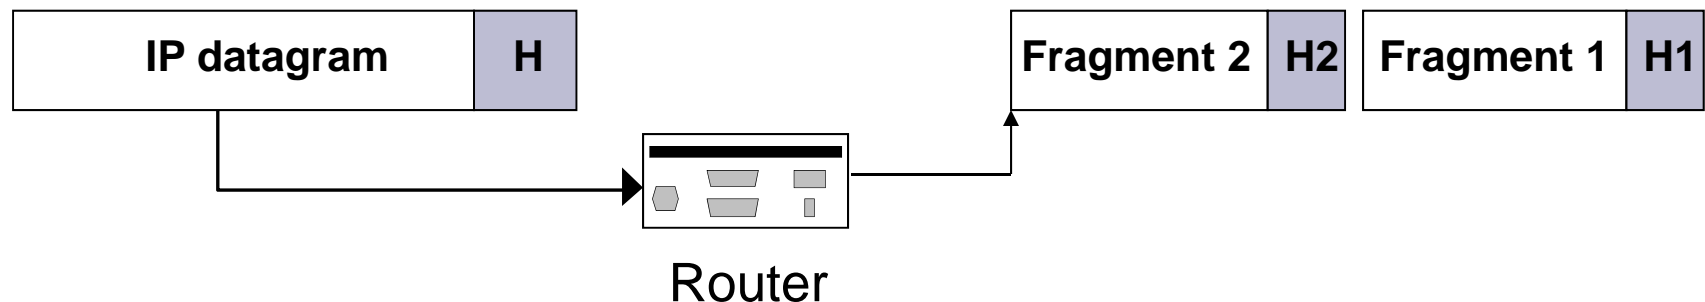
IP Fragmentation

- Host A sends a large IP datagram to Host B.
- How does the intermediate router handle this?
 - IP router splits the datagram into several fragments.
 - Fragmentation requires that the data portion of every fragment except the last be a multiple of 8-bytes.



Where is Fragmentation done?

- Fragmentation can be done at the sender and at intermediate routers.
- The same datagram can be fragmented several times.
- Reassembly of original datagram is only done at destination hosts.




What's involved in Fragmentation?

The following fields in the IP header are involved:

Version	Header Length	Type of Service (TOS)	Total Length (bytes)			
Identification			0	DF	MF	Fragment Offset (8-bytes units)
Time-To-Live (TTL)		Protocol Type	Header Checksum (16 bits)			
...						

- **Identification** is the same in all fragments.
- **Flags** field contains
 - a reserved bit, must be zero,
 - a Don't Fragment (DF) bit that can be set, and
 - a More Fragments (MF) bit.
- **Fragment Offset** contains the offset (in 8-byte units) of current fragment in the original datagram.
- **Total Length** is changed to be the size of the fragment.

Fragmentation through Multiple Links

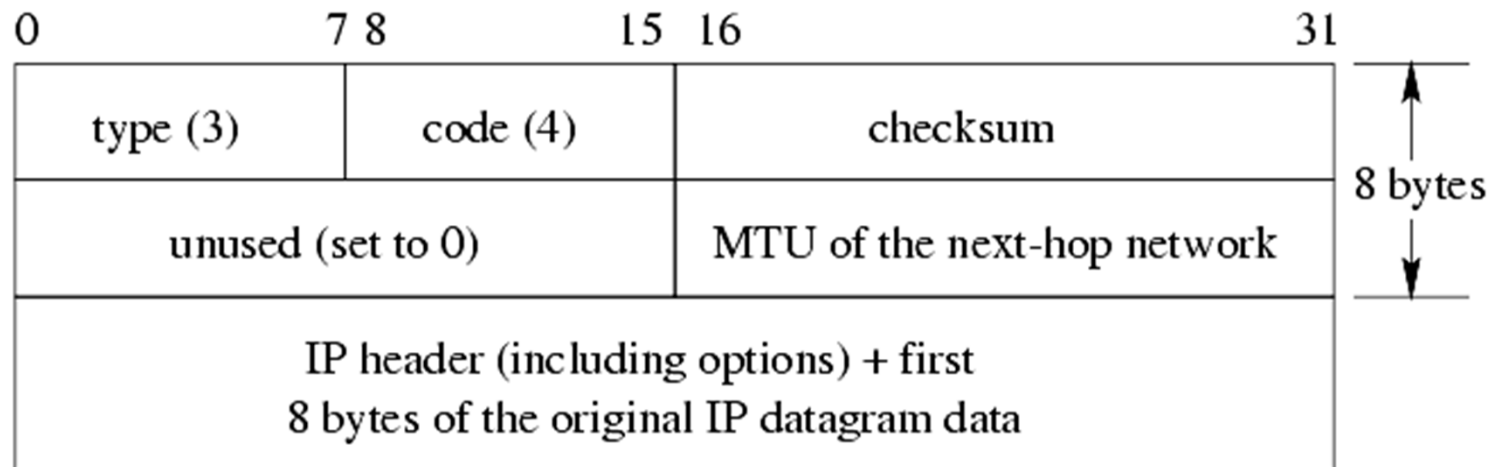


- The ID field stays the same for all fragments of a datagram sent by a sender to allow for reassemble
- The fragment offset is relative to the datagram sent by the sender.
- Two fragments created on X.25 link (offsets 0, 69)
 - $576 - 20$ (IP header) = 556; 552 divides by 8 as 69.
 - First fragment: Offset 0, bytes 1~552; second fragment: Offset 69, bytes 553~1000
- Each fragment is fragmented further on the PPP link
 - ID stays the same on all fragments
 - Fragment offset on the second set of fragments is relative to the original (0, 34, 68, 69, 103)
 - > $296 - 20 = 276$; $272 / 8 = 34$

If the *Don't Fragment* flag is set...


- If fragmentation is needed, and the *Don't Fragment* flag is set, The router drops the datagram and sends an *ICMP unreachable error message* to the source.
- This can be used in *Path MTU Discovery* to find the smallest MTU along a path.

The format of an ICMP unreachable error message:




Path MTU Discovery

A host sends a set of IP datagrams with various lengths and the “don’t fragment” bit set



Transport Layer Protocols

- Transport layer protocols are end-to-end protocols
- They are only implemented at the hosts

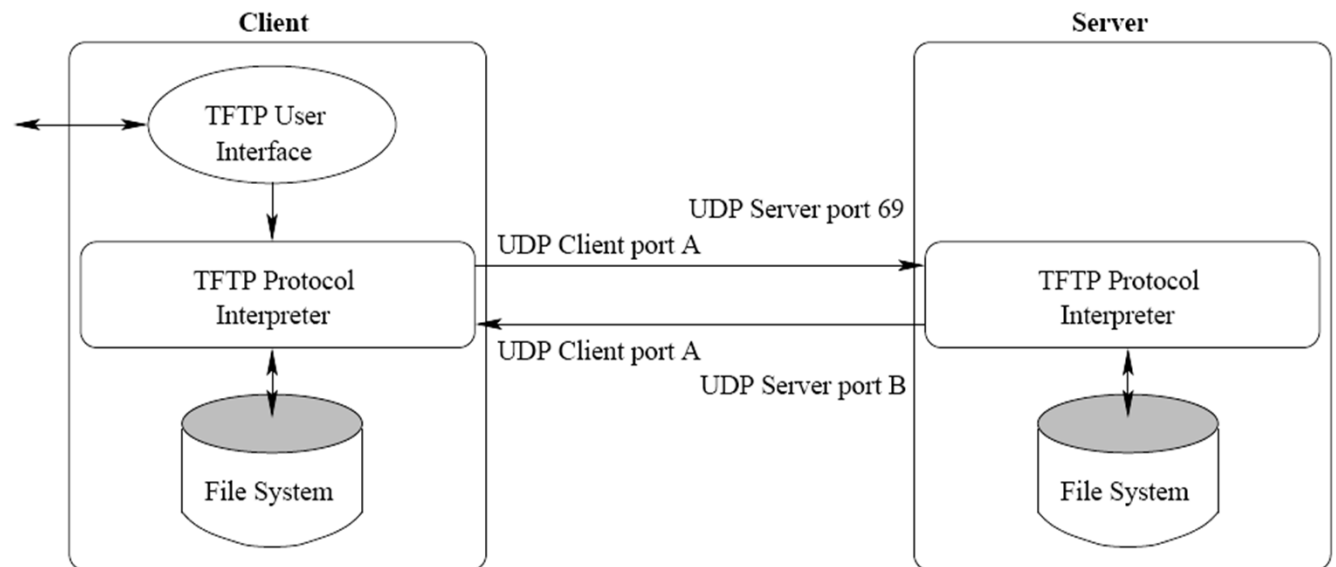


Client-Server Architecture

- Most network applications are implemented using a client-server architecture.
- A server provides network service to the clients.
- Servers use **well-known port numbers** and run these services all the time.
- A client uses a **dynamic port number** and terminates after the service.
- If a client requests a service on a port number not associated with the server,
 - In UDP, an ICMP port unreachable error is returned to the client;
 - In TCP, the TCP connection is reset.

Trivial File Transfer Protocol (TFTP)

- TFTP uses UDP ~ connectionless and unreliable
 - For small infrequent file transfers
 - Throughput is not a major concern
- TFTP uses a stop-and-wait flow window control algorithm
 - Stop for ACK before sending the next data packet
 - A lost packet causes timeout and retransmission
- Designed for diskless systems to download configuration files during bootstrapping



TFTP Packet Format

A typical TFTP session:

1. A client sends a RRQ with a specific filename to a server on UDP port 69
2. If the requested file exists, the server responds with a data packet of length 512 bytes starting with block number 1
3. The client sends an ACK for block number 1
4. The server sends the next data packet with the block number 2
5. The client sends an ACK for block number 2
6. The above two steps continue until the last data block that is shorter than 512 bytes is sent and ACKed

RRQ: Read Request

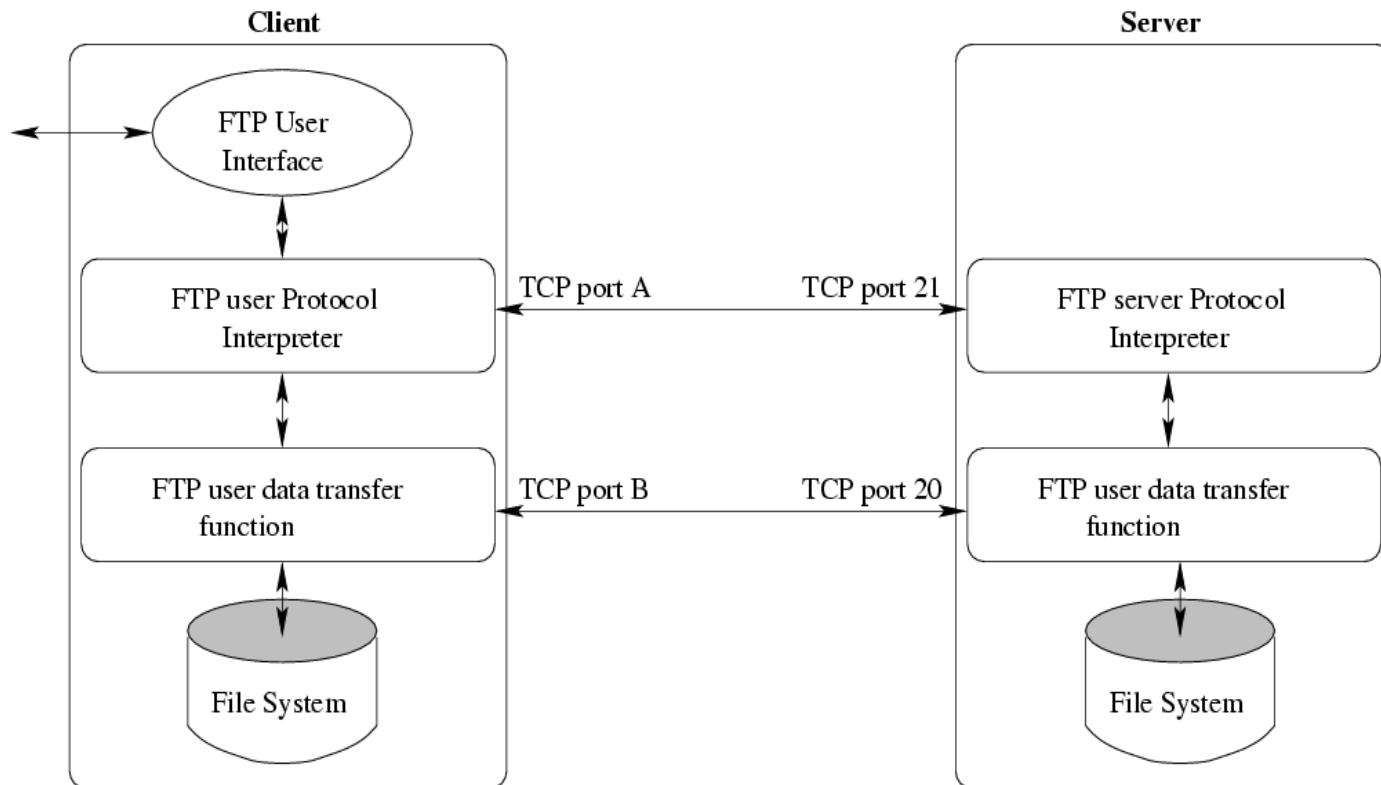
WRQ: Write Request

opcode (1=RRQ, 2=WRQ)	filename	0	mode	0
2 bytes	variable length	1 byte	variable length	1 byte
opcode (3=data)	block number	data		
	2 bytes	0~512 bytes		
opcode (4=ACK)	block number			
opcode (5=error)	block number	error message	0	
		variable length	1 byte	

File Transfer Protocol (FTP)

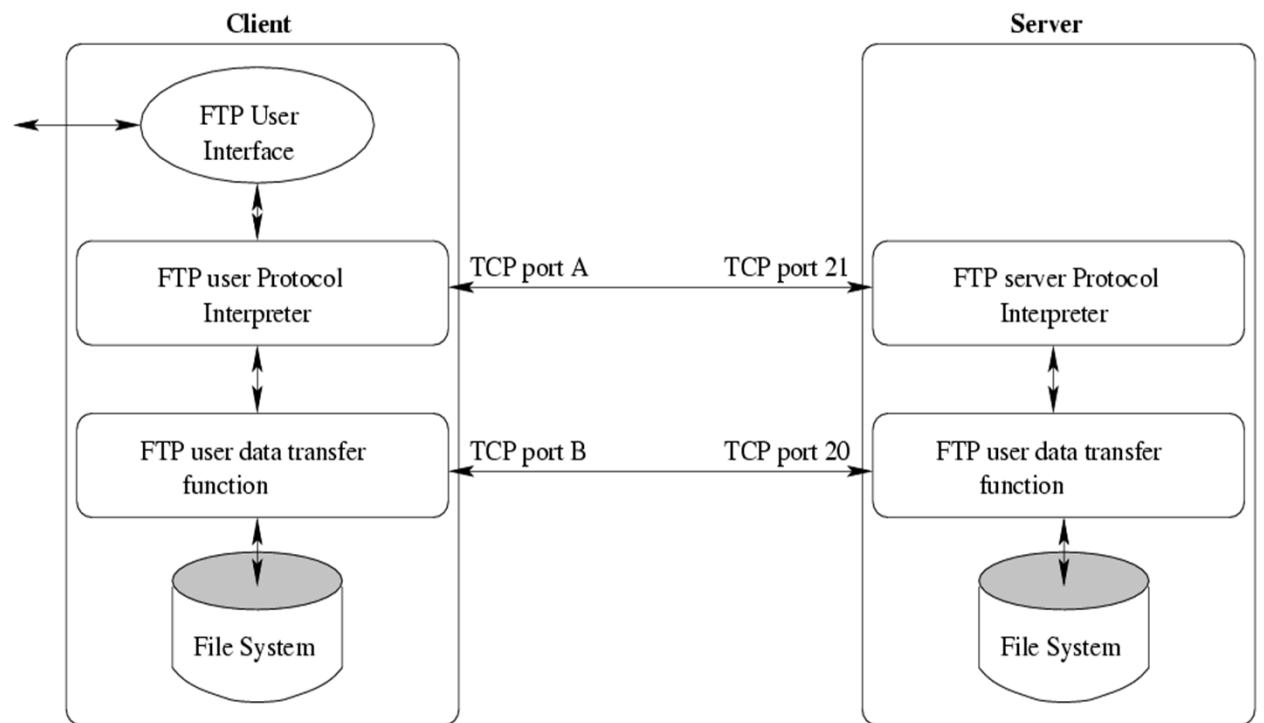
FTP uses two TCP connections

- Control connection: well-known port number at the server = 21
- Data connection: well-known port number at the server = 20



FTP Connection Management

- Control connection, opened by a client, stays up for the duration of the client-server connection.
- Creation of data connection is under the control of a client.
- Client chooses an dynamic port number on the client host for its end of the data connection.
- Client sends the PORT command to the server across the control connection.
- Server receives the port number and issues an active open to that port on the client host. The server uses port 20 at its end for the data connection.
- Multiple FTP sessions from one or more clients to the same FTP server.



FTP Commands



Command	Description
LIST <i>field</i>	list files or directories
PASS <i>password</i>	password on server
PORT <i>n1,n2,n3,n4,n5,n6</i>	client IP address (<i>n1.n2.n3.n4</i>) and port (<i>n5x256+n6</i>)
QUIT	logoff from server
RETR <i>filename</i>	retrieve (get) a file
STOR <i>filename</i>	store (put) a file
TYPE <i>type</i>	specify file type: A for ASCII, I for image
USER <i>username</i>	username on server

FTP Replies



Typical FTP replies

- 125 Data connection already open; transfer starting
- 200 Command OK
- 331 Username OK, password required
- 425 Can't open data connection
- 452 Error writing file
- 500 Syntax error (unrecognized command)
- 501 Syntax error (invalid arguments)

File Transfer: FTP vs. TFTP



FTP

- Complex but reliable file transfer use TCP
- Specified in RFC 959, well-known port 21 (control) and 20 (data)
- Data retransmission carried in lower layer by TCP
- Used for general purpose, high throughput applications
- Security feature provided
 - Username and password checking
 - Data transfer may fail when address translation/firewall implemented with random port passing

TFTP

- Simple and quick file transfer over UDP
- Specified in RFC 1350, well-known UDP port 69 (for originating request to server)
- Both ends use a timeout retransmission to resend a block of data
- Often used to
 - Load into a batch file for multiple hosts
 - Bootstrap diskless systems
- No username and password checking; constitutes a security hole.

Backup Slides



TCP Overview



- Transport layer protocol
- Provides connection-oriented, reliable service to applications, such as HTTP, email, FTP, telnet.
- Only support unicast.
- Features:
 - Error control
 - Flow control
 - Congestion control

TCP Connection

- Source and destination port numbers identify the sending and receiving application processes, respectively.
- **Socket**: the combination of an IP address and a port number.
- A TCP connection is uniquely identified by the two end sockets.

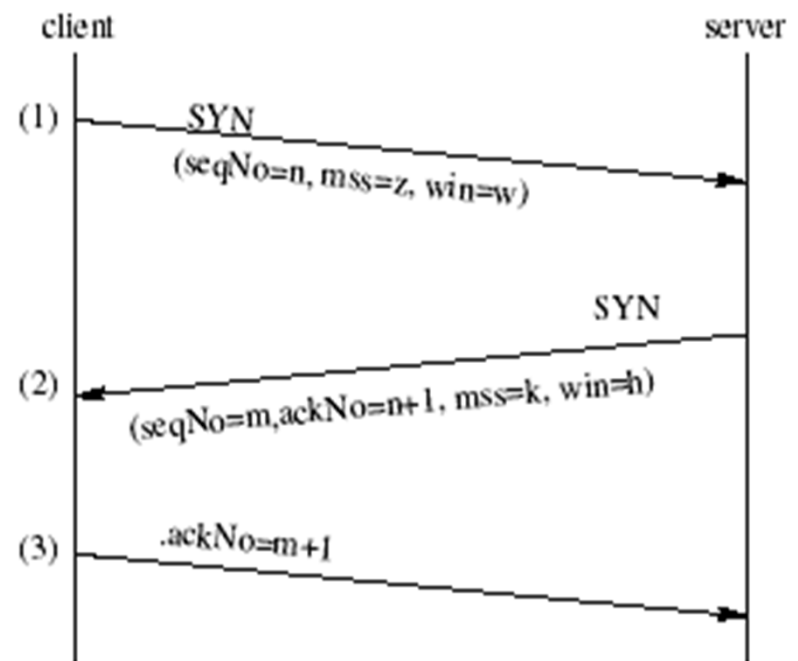
TCP Connection Management

- TCP connection establishment: two end TCP modules
 - allocate required resources for the connection, and
 - Negotiate the value of the parameter uses, such as
 - > Maximum Segment Size (MSS) , given by the receiver side (a.k.a. destination)
 - > Receiving buffer size (i.e. advertised window, WIN), given by the receiver
 - > Initial sequence number (ISN), specified by the sender side (a.k.a. source)
- TCP connection termination

TCP Connection Establishment

Three-way Handshake

- An end host initiates a TCP connection by sending a packet with
 - ISN, say n , in the sequence number field,
 - An empty payload field,
 - MSS,
 - TCP receiving window size, and
 - SYN flag bit is set.
- The other end replies a SYN packet with
 - ACK= $n+1$
 - Its own ISN, say m
 - Its own MSS, and
 - Its own TCP receiving window size
- The initiating host sends an acknowledgement: ACK= $m+1$



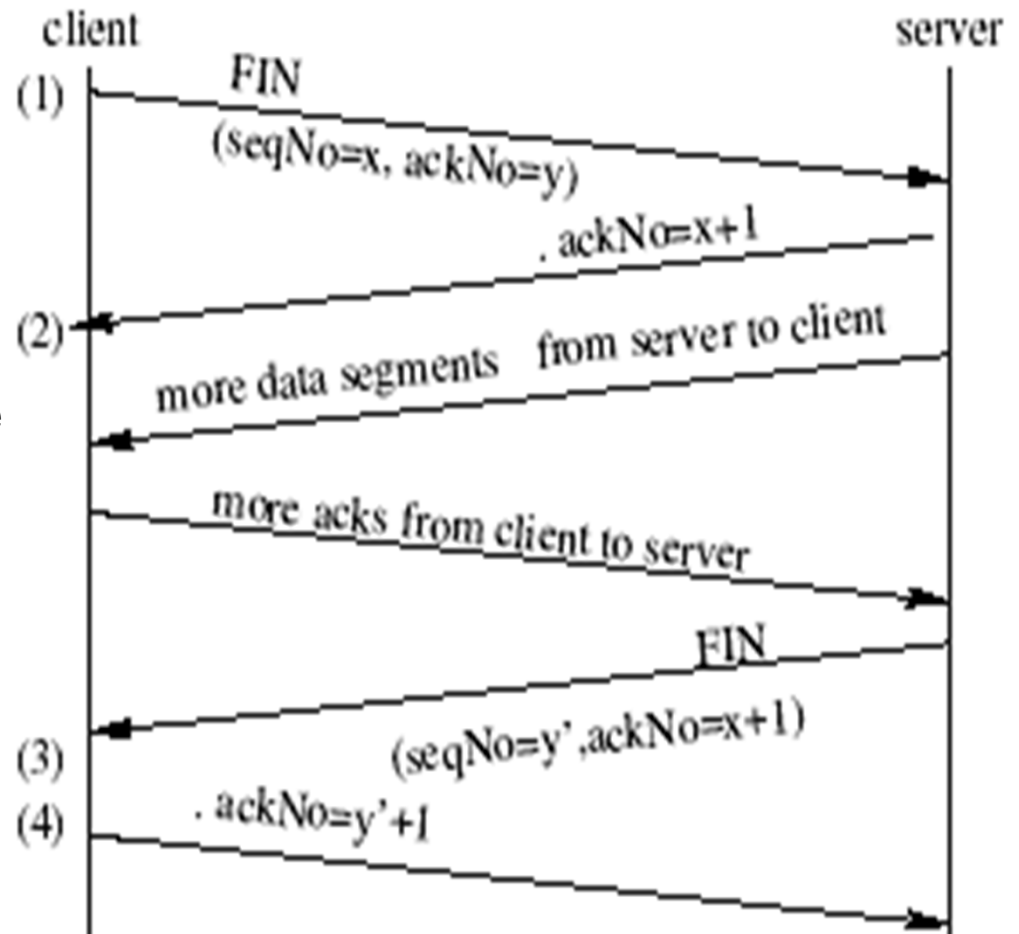
TCP Connection Termination

- A TCP connection is full duplex.
- Each end of the connection has to shut down its one-way data flow towards the other end.
- After termination performed, the connection must stay in the `TIME_WAIT` state for twice the *Maximum Segment Life* (MSL) to wait for delayed segments.
- If an unrecoverable error is detected, either end can close the TCP connection by sending a RST segment.

TCP Connection Termination (cont'd)

Four-way handshake

- TCP Half-Close
 - One end TCP sends a packet with the FIN flag set.
 - The other end acknowledges the FIN segment.
 - The data flow in the opposite direction still works.
- Do another Half-Close in the opposite direction.



TCP Data Flow



- TCP provides a byte-stream connection to the application layer.
- The sender TCP module
 - Receives a byte stream from the application and puts the bytes in a sending buffer.
 - Extracts the bytes from the sending buffer and sends to the lower network layer in blocks (TCP segments).
- The receiver TCP module
 - Uses a receiving buffer to store and reorder received TCP segments.
 - Restores a byte stream from the receiving buffer and sends to the application process.